

Option 1: Polyconvexity meets learned anisotropy:
Paired directional features for hyperelastic neural networks

Option 2: Learning anisotropy with convex neural networks:
Paired directional features for polyconvex hyperelasticity

S. Ares de Parga^{1,4,†} A. Cornejo^{1,2} J. A. Hernández^{1,3} R. Rossi^{1,2}

¹Centre Internacional de Mètodes Numèrics en Enginyeria (CIMNE), Barcelona, Spain

²Universitat Politècnica de Catalunya, Department of Civil and Environmental Engineering (DECA), Barcelona, Spain

³Universitat Politècnica de Catalunya, Departament de Resistència de Materials i Estructures a l'Enginyeria, E.S. d'Enginyeries Industrial, Aeroespacial i Audiovisual de Terrassa (ESEIAAT), Terrassa, Spain

⁴Department of Aeronautics and Astronautics, Stanford University, Stanford, CA 94305, USA

[†] This work was conducted while S. Ares de Parga was affiliated with CIMNE.

10 September 2026 — Working draft v0.2

Abstract

Approximating the effective anisotropic hyperelastic response of a representative volume element (RVE) requires constitutive expressiveness, mechanical consistency, and computational efficiency. Non-intrusive neural surrogates provide direct constitutive evaluations, whereas intrusive projection-based reduced-order models (PROMs) approximate the microscopic state and equilibrium equations. This work develops a polyconvex physics-augmented neural network (PANN) construction based on paired directional features with trainable orientations, stretch powers, and determinant exponents. The features are convex in independent deformation minors and are composed with convex, componentwise nondecreasing cores realized as an input-convex neural network (ICNN) or an input-convex Kolmogorov–Arnold network (ICKAN). Their isotropic reference derivatives enable affine-determinant stress normalization without compromising polyconvexity. The complete energies satisfy objectivity, an exactly stress-free reference, two-dimensional polyconvexity, and prescribed growth conditions by construction; global nonnegative energy is assessed through a separate sufficient bound for the trained parameters. Projection-based surrogates with hyperreduction provide a complementary benchmark for structural accuracy and online cost. On a periodic cell with a rotated elliptical void, the selected polyconvex models attain approximately 0.65% held-out stress error. A tensile-coupon application illustrates structural deployment against a full-order FE² reference. The ICNN and ICKAN give displacement errors below 0.43%; the iterative nonlinear reduced model gives 0.052% error and its implemented workflow is 2.23 times faster than the optimized affine reduced model under the reported configurations. These results demonstrate the practical use of the constrained constitutive representation and quantify its accuracy–cost trade-off against reduced microscopic solvers.

Keywords: computational homogenization; hyperelasticity; physics-augmented neural networks; polyconvexity; projection-based reduced-order models; empirical cubature.

1 Introduction

The effective response of a heterogeneous material emerges from the deformation of its constituents and their interactions with the microstructure. At finite strains, this interaction can produce anisotropy and coupling between loading modes that are difficult to capture with a predetermined macroscopic constitutive law. Computational homogenization addresses this difficulty by resolving a representative volume element (RVE): an imposed deformation drives a microscopic boundary-value problem, from which effective stresses and constitutive derivatives are obtained [1–3]. Rather than prescribing the effective law in advance, the RVE supplies it through the solution of the underlying mechanical problem.

This flexibility comes at a computational cost. Characterizing the material over a range of deformations already requires many nonlinear finite element solves; using the same RVE within a structural calculation multiplies those solves across macroscopic integration points and load increments [3, 4]. Much of this repeated work concerns closely related microscopic states. Surrogate modeling seeks to exploit that relationship by

computing representative solutions offline and reusing their information in less expensive online evaluations. The objective is therefore not only to reduce execution time and the resources required to store and process microscopic fields, but to do so without losing the constitutive behavior that made the RVE description useful.

Two complementary approximation routes are relevant, but they play different roles in this work. Non-intrusive constitutive surrogates learn the effective response from RVE input–output data and dispense with the microscopic solve online [5]. Intrusive projection-based reduced-order models (PROMs) instead compress the microscopic displacement field and retain projected equilibrium equations [6, 19, 20], with hyperreduction limiting the local contributions required to evaluate them [34]. We focus on the first route: learning the effective stored energy while embedding the constitutive restrictions needed for its mechanical use. The intrusive route is retained as a secondary reference because it preserves an approximate microscopic state and the underlying constituent law.

The central contribution is a family of learnable paired directional features for polyconvex anisotropic energies. Their orientations and exponents adapt to the RVE response, while their paired construction permits an admissible exact normalization at the undeformed state. Coupling these features to either an input-convex neural network (ICNN) or an input-convex Kolmogorov–Arnold network (ICKAN) gives two realizations of the same constitutive design. PROM-based surrogates of the same RVE then provide secondary accuracy–cost references, and a tensile structural example tests how the learned energies behave beyond isolated material evaluations.

1.1 Energy-based constitutive learning and polyconvexity

Direct constitutive learning moves the expensive microscopic solves to the data-generation stage and replaces them online by a trained material model. Neural homogenization has demonstrated this separation [7], making systematic sampling of representative microscopic responses an essential part of the offline workflow [8]. What the surrogate must represent depends on the material class: history-dependent responses may require learned history-to-response maps [9] or internal descriptors such as maximum directional stretches [10], while interacting scales may motivate staggered data generation and training [11]. Here, by contrast, we consider history-independent hyperelasticity. This removes memory from the state but imposes a precise constitutive requirement: the stress must derive from a stored-energy function.

Learning that function, rather than fitting each stress component independently, makes the stress work conservative and gives a major-symmetric constitutive tangent by differentiation. This is a central motivation for the mechanics-informed energy formulation of As’ad et al. [12]. Expressing the energy through objective deformation measures then removes dependence on a superposed rigid rotation. In this way, physics-augmented neural networks (PANNs) shift part of the burden from learning mechanical identities from data to satisfying them through the representation itself [13, 14]. Prediction accuracy remains to be established, but potential structure and objectivity need not depend on how densely the training domain has been sampled.

Preserving a potential does not, by itself, control its curvature. At finite strains, that control must also be imposed in appropriate variables. Convexity in Green–Lagrange strain, as used in Ref. [12], is attractive because stress and tangent follow directly from the learned energy. Its implication for curvature with respect to the deformation gradient nevertheless contains a stress-dependent term [15], so it is not a global polyconvexity guarantee. Conversely, Xu et al. [16] parameterize an incremental tangent through a lower-triangular factorization to enforce positive semidefiniteness. This controls the tangent matrix but does not by itself define the scalar potential required for integrability. The modeling task is therefore to preserve both the potential and the desired finite-strain curvature condition within one construction.

Polyconvexity provides a route to that combination. It represents the energy as a convex function of deformation minors, yielding a sufficient condition for rank-one convexity [17]. For the two-dimensional model considered here, the independent arguments are the deformation gradient and its determinant. This framework also underpins existence results when the additional coercivity and admissibility hypotheses are met [18]. We use it as a sufficient constitutive restriction, not as a necessary property of every homogenized material or a guarantee of structural uniqueness. The remaining challenge is representational: the restrictions that provide a guarantee must still leave enough freedom to describe the anisotropic RVE response.

1.2 Directional features and representational expressiveness

The remaining representational challenge is especially consequential for anisotropy: a flexible scalar core can distinguish only the deformation states that its inputs distinguish. Structural tensors encode material directions in objective scalar quantities and underpin established anisotropic polyconvex energies [40]. The directional description is itself a constitutive choice. Models of fiber-reinforced tissue, for example, must account for a distribution of orientations rather than one perfectly aligned family when that dispersion affects

the mechanical response [41]. A learned feature map must likewise expose the relevant directional distinctions while remaining compatible with objectivity and convexity in the deformation minors. It is therefore part of the constitutive model, not merely a preprocessing step.

A convex scalar core does not resolve this choice. An ICNN provides a convex function [42], but a sufficient composition with nonlinear features requires those features to be convex in the independent deformation minors and the core to be componentwise nondecreasing [13]. This monotonicity imposes an order on feature space: if $\mathbf{z}(\mathbf{F}_a) \leq \mathbf{z}(\mathbf{F}_b)$ componentwise, then $H(\mathbf{z}(\mathbf{F}_a)) \leq H(\mathbf{z}(\mathbf{F}_b))$. Increasing the network width cannot reverse that ordering. When reference corrections and growth terms complete the energy, the comparison applies to the contribution represented by the core, but the restriction remains. A poorly chosen feature family can therefore impose a representational obstruction before network capacity or optimization becomes limiting.

The numerical evidence of Klein et al. [13] illustrates this distinction. For their cubic beam-lattice data, an exactly objective invariant-based polyconvex model fails to represent the shear response, whereas a model formulated in $\mathbf{F}$, $\text{cof } \mathbf{F}$, and $\det \mathbf{F}$ gives a substantially better fit while approximating objectivity through data augmentation. The same invariant-based model performs very well for their analytically generated transversely isotropic data. These results do not identify polyconvexity alone as the source of limited accuracy; they show that success depends on whether the admissible inputs expose the distinctions required by the material response. This motivates enlarging the feature family before concluding that a convex core is inadequate.

The natural response is to make the directional description adaptive without abandoning the admissible composition. Generalized-structure-tensor models calibrate the structural tensors together with the energy model, allowing the data to identify aspects of the anisotropy [43]. Recent triclinic invariant constructions and group symmetrization further broaden the polyconvex descriptions available for finite material-symmetry groups [44]. Greater directional freedom, however, must be reconciled with the undeformed state. A learned anisotropic contribution need not produce zero reference stress, while cancelling that stress with an arbitrary strain correction can violate polyconvexity or material symmetry. The invariant-based normalization and analytic growth terms of Linden et al. [14] demonstrate why these operations must belong to the admissible constitutive construction. The feature family must therefore do two jobs at once: distinguish the relevant directional response and control its reference derivatives so that exact stress normalization remains compatible with polyconvexity. This is the role of the paired features introduced here.

The feature map determines what the model can distinguish; the scalar core determines how those distinctions are combined into an energy. Prior constructions distribute this flexibility in different ways. Constitutive artificial neural networks (CANNs) learn coefficients and exponents within an interpretable catalog of admissible energy terms [45], whereas neural ordinary differential equations can learn monotone energy derivatives and recover convex energy contributions by integration [46]. In the present work, an ICNN combines the paired features through nonnegative weights and fixed convex activations. A Kolmogorov–Arnold network (KAN) instead places trainable univariate functions on its edges, a flexibility explored in constitutive KANs [47]. A scalar energy output alone does not preserve the feature-composition argument; the edge functions must satisfy the required convexity and monotonicity, as in monotone input-convex KANs (ICKANs) [48]. We therefore compare a smooth ICNN with an integrated-hat ICKAN under the same paired features and analytic completion, joining the learned splines to admissible tails. The comparison isolates how the scalar dependence is represented rather than changing the directional information or the mechanical guarantees.

1.3 Projection-based reduced-order models

The preceding approach learns the effective energy and thereby removes the microscopic solve from the online stage. Projection-based reduction makes a different compromise: it retains the RVE boundary-value problem, but seeks its solution in a space learned from previous microscopic states. In a conventional PROM, those states define a basis, usually through proper orthogonal decomposition (POD), and projected equilibrium determines the coordinates of each new displacement field. Reduced finite-strain homogenization has shown that this strategy can reuse information from earlier RVE solves without replacing the constituent law [19, 20]. Moreover, because the result is still an approximate microscopic field, it can support settings in which reduced periodic substructures exchange kinematic information [21]. PROMs therefore provide a natural physics-based reference for the present study: they preserve more of the microscopic problem than an effective-energy surrogate, but pay for it by retaining an online equilibrium solve.

The affordability of that solve depends first on the geometry of the solution family. A small number of loading parameters does not necessarily imply that the corresponding microscopic states lie close to a small global linear space. When their Kolmogorov n -width decays slowly, many basis vectors—and hence many equilibrium unknowns—are required to achieve adequate accuracy. Considerable effort has therefore been

devoted to mitigating this Kolmogorov barrier with richer reduced representations [22–27]. Local methods construct a piecewise-linear approximation by replacing one global space with several spaces associated with different regions of the solution set [25]. Manifold methods go further by allowing a few coordinates to describe a curved approximation, using constructions such as autoencoder mappings [26], quadratic corrections [27], or learned closures within a POD coordinate space [22–24]. The last of these offers the particular compromise sought here: PROM–ANN enriches the geometry of the reduced state without giving up either its POD description or projected equilibrium [22]; a subsequent hypersonic-flow assessment compared this nonlinear representation with affine and quadratic alternatives [50]. Equilibrium determines only a small set of primary coordinates, while an artificial neural network (ANN) reconstructs the secondary coordinates that populate a much richer POD expansion.¹ The microscopic displacement can thereby draw on many modes while the online nonlinear system remains small, which motivates the primary–secondary construction adopted for our RVE.

Reducing the number of unknowns, however, does not by itself make their equilibrium equations inexpensive. At every nonlinear iteration, assembling the reduced residual and tangent may still require constitutive evaluations over the complete microscopic mesh. Faced with this remaining cost, one may either perform the same full-mesh work faster or reduce how much of it must be performed. Hardware acceleration pursues the first option [28]. Hyperreduction pursues the second, reconstructing the required operators from sparse information and thereby turning a PROM into a hyperreduced PROM (HPROM). Choosing this latter path raises a further design question: should the high-dimensional nonlinear quantities be approximated before projection, or should projection come first so that only the reduced equations are approximated? These alternatives define the approximate-then-project and project-then-approximate viewpoints, respectively [29, 30]. The former reconstructs a sampled high-dimensional quantity before projecting it, as exemplified by discrete empirical interpolation (DEIM) [31]. The latter acts directly on the reduced equations, replacing full assembly by weighted contributions from a small subset of finite elements or other spatial entities. This route includes the linear-programming empirical quadrature procedure (EQP) for reduced-basis models [35, 36], energy-conserving sampling and weighting (ECSW) [32, 33], and the empirical cubature method (ECM) [21, 34]. The structure-preservation results obtained for ECSW in structural dynamics reinforce the value of targeting the projected equations [33], although they do not extend automatically to every sampling-and-weighting construction. Within the elementwise finite element setting adopted here, both ECSW and ECM are natural choices: each replaces full assembly by a nonnegative weighted sum of selected contributions while leaving the local constituent evaluations intact. We follow ECM because a sequence of developments in our previous work extends its fixed rule toward movable points and subspace- and manifold-adaptive weights [37–39], providing the methodological path needed to couple hyperreduction with PROM–ANN.

That path begins with a limitation of the classical rule. ECM can remove most element evaluations, but one fixed support and one vector of nonnegative weights must represent the entire training space [34]. CECM first relaxed the spatial part of this restriction: starting from a feasible discrete rule, it readjusts both point positions and weights during construction so that the final support is no longer confined to the original candidates [37]. This additional freedom can produce a sparser rule, but the resulting weights must still serve the full state space. SAW–ECM addressed that remaining globality for a collection of local linear spaces. It retains a common support, yet assigns a different nonnegative weight vector to each subspace, allowing the rule to follow which local representation is active [38]. Once those discrete local spaces are replaced by a continuously parameterized nonlinear manifold, the natural continuation is to make the weights continuous as well. MAW–ECM therefore promotes the subspace-specific vectors to weight fields over the latent coordinates [39]. This final step closes the loop with PROM–ANN: the same primary coordinates that locate the microscopic state on the learned manifold can drive both the decoder of its secondary coordinates and the adaptive weights used to evaluate its equilibrium and stress.

This sequence produces a controlled ladder of intrusive references for the structural comparison. The affine HPROM retains a linear POD state, solves projected equilibrium, and evaluates it with fixed ECM. HPROM–ANN replaces the linear state by the primary–secondary manifold and the fixed rule by adaptive integration, but still determines its primary coordinates through equilibrium. D–HPROM–ANN removes that final online solve by inferring the microscopic state directly from the imposed strain. Moving across the three models therefore removes successive portions of the online microscopic work while preserving an approximate microscopic state and selected evaluations of the constituent law. Their purpose is to contextualize the accuracy and cost of the proposed energy surrogates, not to supply the same constitutive guarantees: neither the retained constituent law nor nonnegative cubature weights alone ensures that the approximate effective stress is the derivative of a single effective potential.

¹The closure need not be neural: the latent-space framework of Ares De Parga et al. [24] also considers Gaussian-process regression (GPR) and radial-basis-function (RBF) closures.

1.4 Present contribution

Against these intrusive physics-based references, the central route pursued in this work is an effective-energy surrogate whose anisotropy is learned without relinquishing polyconvexity or exact reference-state normalization. These requirements are coupled: a learnable anisotropic contribution generally produces a nonzero stress at the undeformed state, whereas cancelling that stress with an arbitrary correction can destroy the admissible constitutive structure. We resolve this tension through paired directional features with trainable orientations, stretch powers, and determinant exponents constrained by explicit joint-convexity bounds.

The pairing is designed so that the strain derivative of every feature at the undeformed state is proportional to the identity tensor, even though its response away from that state remains anisotropic. The complete learned contribution can consequently be normalized by a single scalar correction affine in the determinant. This correction cancels the reference stress without tying that cancellation to the independently adjustable volumetric stiffness, allowing the directional description to adapt to the RVE while the reference state remains under analytic control.

Once the directional information has been exposed in this admissible form, two monotone convex cores provide alternative representations of its scalar energetic dependence: a smooth ICNN and an integrated-hat spline ICKAN with analytically matched tails. In both cases, the learned core, reference correction, and growth terms form one complete energy for which the analysis establishes objectivity, an exactly stress-free reference, two-dimensional polyconvexity, and the specified growth limits. Global nonnegative energy is treated separately through a sufficient lower bound evaluated for the trained parameters. The novelty lies in bringing established convex-network and affine-determinant normalization principles together through the paired anisotropic family and its explicit analytical conditions.

Figure 1 provides a graphical synopsis of this construction, whose individual blocks are developed in the sections that follow. It traces the path from RVE data through the paired directional features and either convex core to the analytic completion of the stored energy, whose derivatives supply stress and tangent. The numerical study then places these models beside a flexible energy reference to interpret the cost of the guarantees and beside the reduced microscopic evaluators just described to contextualize structural accuracy and online cost.

1.5 Assessment strategy and scope

The assessment follows the two levels at which the surrogates are intended to operate. At the constitutive level, it asks how accurately the constrained energies represent independent anisotropic RVE responses and what is gained by learning the paired directions rather than prescribing them. Response errors and the fixed/learned-feature comparison address approximation quality, while derivative checks and saved-parameter energy bounds interrogate behavior that a fit metric alone cannot reveal [49]. Sampled curvature tests serve here as implementation diagnostics; the polyconvexity guarantee itself comes from the analytical construction.

At the structural level, the question becomes whether those constitutive approximations remain accurate and useful after being embedded in a nonlinear boundary-value problem. A tensile coupon provides this deployment test. Within the first-order two-level finite element formulation, denoted FE^2 , an RVE problem supplies the material response at each macroscopic integration point [4]; here every point queries the same periodic cell. Replacing that query by a surrogate propagates its approximation through a common equilibrium problem. A nested full-order model (FOM)– FE^2 calculation supplies the reference, while field-level comparisons complement aggregate displacement and stress errors. The intrusive reduced models then provide physics-based accuracy–cost benchmarks without turning the development of a new FE^2 scheme into a second objective of the paper.

These two levels deliberately separate analytical guarantees from empirical adequacy. The constitutive construction establishes the stated properties of the complete energy, whereas independent RVE and structural evaluations assess how well that admissible representation approximates the selected material and how much online work it avoids. Additional closed-cycle and curvature diagnostics are documented in the supplementary material.

The resulting claims are necessarily tied to the scope of the study: material A, plane-strain kinematics, and the loading path applied to the coupon. Agreement with FOM– FE^2 measures surrogate error within the adopted homogenization model, not error relative to a structure in which every microscopic cell is resolved explicitly. The comparisons assess the feature family, normalization, and convex core as a combined constitutive design rather than attributing performance exhaustively to each component. Reported timings concern the implemented online workflows; offline data generation, training, memory savings, and break-even behavior remain outside the comparison. The objective within these boundaries is to establish the usefulness and constitutive properties of the proposed representation, not universality across microstructures, general




Figure 1: Overview of the proposed constitutive construction. Sampled strain states are evaluated with the periodic RVE model to supply energy and stress data. These responses train learnable paired directional features and a monotone convex ICNN or ICKAN core. Analytic reference and growth terms complete the stored energy; differentiation then supplies the work-conjugate stress and material tangent. The stated properties follow from the complete construction rather than from sampled numerical checks.

three-dimensional completeness, or robustness under every possible microscopic instability.

With this assessment strategy established, Section 2 defines the homogenization problem and work-conjugate variables. Section 3 introduces the admissibility criteria, constructs the learned energies, and proves their properties, after which Section 4 defines the reduced microscopic references used in the structural comparison. Sections 5 and 6 report the constitutive and structural assessments, respectively; Section 7 closes the paper.

2 Finite-strain homogenization and the reference problem

The two surrogate routes replace different parts of the online calculation, but both approximate the same constitutive map: for a prescribed macroscopic deformation, determine the effective response of a periodic RVE. The unreduced RVE model supplies the reference by solving microscopic equilibrium. This section establishes the kinematic input, periodic cell problem, effective outputs, and stress–tangent interface shared by all models before their approximation strategies diverge. The cell geometry and material parameters are specified in the numerical examples.

2.1 Kinematics and work-conjugate variables

Let $\mathbf{F} \in \text{GL}^+(2)$ denote the prescribed in-plane macroscopic deformation gradient. Its associated right Cauchy–Green tensor and Green–Lagrange strain are

$$\mathbf{C} = \mathbf{F}^T \mathbf{F}, \quad \mathbf{E} = \frac{1}{2}(\mathbf{C} - \mathbf{I}), \quad J = \det \mathbf{F} > 0. \quad (1)$$

Here $\mathbf{C}$ and $\mathbf{E}$ describe the macroscopic state. The macroscopic input is written simply as $\mathbf{F}$ to match the constitutive notation used throughout the paper. The subscript μ marks microscopic quantities, whereas the unsubscripted $\mathbf{W}$, $\mathbf{P}$, and $\mathbf{S}$ denote effective quantities obtained by homogenizing the cell. We adopt plane strain: the corresponding three-dimensional gradient is $\text{diag}(\mathbf{F}, 1)$, and the same embedding applies locally to $\mathbf{F}_\mu$. This constrains the out-of-plane deformation, not the out-of-plane stress.

The constitutive input used throughout is the three-component engineering representation of the Green–Lagrange strain. With $\mathbf{S}$ the effective second Piola stress, the work-conjugate convention is

$$\mathbf{e} = (E_{11}, E_{22}, 2E_{12})^T, \quad \mathbf{s} = (S_{11}, S_{22}, S_{12})^T, \quad \mathbf{S} : d\mathbf{E} = \mathbf{s} \cdot d\mathbf{e}. \quad (2)$$

Thus $\gamma_{12} = 2E_{12}$ is paired with S_{12} , with no additional factor of two in the stress vector. For a differentiable effective energy density $W(\mathbf{e})$, this convention gives $\mathbf{s} = \partial_{\mathbf{e}} W$ and $\mathbf{D} = \partial_{\mathbf{e}} \mathbf{s}$. It is used consistently in training, work integrals, and tangent checks.

An admissible strain input must satisfy $\mathbf{I} + 2\mathbf{E} \succ 0$. We query the RVE with the representative

$$\mathbf{F}(\mathbf{e}) = \sqrt{\mathbf{I} + 2\mathbf{E}}, \quad (3)$$

where the square root is symmetric positive definite. Objectivity allows a superposed macroscopic rotation to be removed without changing the energy or second Piola stress. This choice does not assume isotropy: material directions remain fixed in the reference configuration. Nor does it require the local microscopic gradients to be symmetric.

2.2 Periodic microscopic equilibrium

Let Y denote the reference periodic cell and $Y_s \subset Y$ its solid part. The prescribed $\mathbf{F}$ determines the affine part of the microscopic deformation; the fluctuation $\mathbf{w}$ is unknown:

$$\boldsymbol{\varphi}(\mathbf{X}) = \mathbf{F}\mathbf{X} + \mathbf{w}(\mathbf{X}), \quad \mathbf{F}_\mu = \mathbf{F} + \nabla \mathbf{w}, \quad (4)$$

with gradients taken in the reference configuration. Periodicity requires $\mathbf{w}(\mathbf{X}^+) = \mathbf{w}(\mathbf{X}^-)$ at paired points on opposite cell faces. Consequently, the microscopic displacement $\mathbf{u}_\mu = \boldsymbol{\varphi} - \mathbf{X}$ satisfies

$$\mathbf{u}_\mu(\mathbf{X}^+) - \mathbf{u}_\mu(\mathbf{X}^-) = (\mathbf{F} - \mathbf{I})(\mathbf{X}^+ - \mathbf{X}^-). \quad (5)$$

The finite-strain displacement jump therefore depends on $\mathbf{F} - \mathbf{I}$, not directly on $\mathbf{E}$. Setting one fluctuation value to zero removes the arbitrary translation; it does not clamp the entire cell boundary. The internal void surface is traction-free.

Let $\psi_\mu(\mathbf{X}, \mathbf{F}_\mu)$ be the solid's objective hyperelastic energy per unit reference volume and $\mathbf{P}_\mu = \partial_{\mathbf{F}_\mu} \psi_\mu$ its first Piola stress. In the absence of microscopic body forces, equilibrium requires

$$\int_{Y_s} \mathbf{P}_\mu : \nabla \delta \mathbf{w} \, dV = 0 \quad (6)$$

for every admissible periodic test fluctuation. Periodic kinematics and the associated traction balance provide the standard work-consistent cell problem [1]. After finite element discretization and elimination of the periodic constraints, let $\mathbf{a}$ collect the independent fluctuation coordinates. The full-order equilibrium problem is

$$\mathbf{r}(\mathbf{a}, \mathbf{e}) = \partial_{\mathbf{a}} \Pi_\mu(\mathbf{a}, \mathbf{e}) = \mathbf{0}, \quad \Pi_\mu = \int_{Y_s} \psi_\mu(\mathbf{X}, \mathbf{F}_\mu) \, dV, \quad (7)$$

where the strain dependence enters through Eq. (3). Solving Eq. (7) gives $\mathbf{a}^*(\mathbf{e})$ and hence the microscopic fields. The microscopic deformation must also remain orientation-preserving, $\det \mathbf{F}_\mu > 0$.

2.3 Effective energy and stress

On the selected equilibrium branch, the effective energy and stresses are

$$W(\mathbf{e}) = \frac{\Pi_\mu(\mathbf{a}^*(\mathbf{e}), \mathbf{e})}{|Y|}, \quad \mathbf{P} = \frac{1}{|Y|} \int_{Y_s} \mathbf{P}_\mu \, dV, \quad \mathbf{S} = \mathbf{F}^{-1} \mathbf{P}. \quad (8)$$

The normalizing volume $|Y|$ includes the void: these are properties of the porous cell, not of its solid fraction alone. For the two-dimensional discretization, the same reference thickness enters the integration weights and the cell volume, and cancels in their ratio. The stress conversion in Eq. (8) is performed *after* averaging the first Piola stress. In general, averaging microscopic second Piola stresses does not produce $\mathbf{S}$, because $\mathbf{F}_\mu$ varies within the cell.

The connection between energy and averaged stress follows from equilibrium. For a macroscopic variation, $\delta \mathbf{F}_\mu = \delta \mathbf{F} + \nabla \delta \mathbf{w}$; Eq. (6) removes the fluctuation contribution and yields

$$\frac{1}{|Y|} \int_{Y_s} \mathbf{P}_\mu : \delta \mathbf{F}_\mu \, dV = \mathbf{P} : \delta \mathbf{F} = \mathbf{S} : \delta \mathbf{E} = \mathbf{s} \cdot \delta \mathbf{e}. \quad (9)$$

This is the Hill–Mandel work-consistency condition [1]; it also establishes $\mathbf{s} = \partial_{\mathbf{e}} W$ on a smooth branch.

The Newton solver follows an equilibrium branch, possibly through load continuation. Convergence establishes stationarity, not global minimization of Π_μ or uniqueness of the cell response. We use the resulting smooth branch as the reference constitutive response over the investigated states; we do not assume that microscopic polyconvexity guarantees a polyconvex homogenized energy.

2.4 Consistent constitutive tangent

The tangent must account for the change in microscopic equilibrium as the input strain changes. To make this dependence explicit, write the stress evaluator before substitution of the equilibrium state as $\widehat{\mathbf{s}}(\mathbf{a}, \mathbf{e})$, so that $\mathbf{s}(\mathbf{e}) = \widehat{\mathbf{s}}(\mathbf{a}^*(\mathbf{e}), \mathbf{e})$. If $\mathbf{r}_{,\mathbf{a}}$ is nonsingular on the selected branch, implicit differentiation of Eq. (7) gives

$$\frac{d\mathbf{a}^*}{d\mathbf{e}} = -\mathbf{r}_{,\mathbf{a}}^{-1} \mathbf{r}_{,\mathbf{e}}, \quad \mathbf{D} = \widehat{\mathbf{s}}_{,\mathbf{e}} - \widehat{\mathbf{s}}_{,\mathbf{a}} \mathbf{r}_{,\mathbf{a}}^{-1} \mathbf{r}_{,\mathbf{e}}. \quad (10)$$

All partial derivatives are evaluated at equilibrium, holding the other argument fixed. The first term differentiates the stress at fixed fluctuation coordinates; the second accounts for microscopic relaxation. Omitting that correction generally does not give the derivative of the deployed constitutive response. For the smooth hyperelastic branch, the result is also $\partial_{\mathbf{e}\mathbf{e}}^2 W$.

No explicit matrix inverse is needed in Eq. (10): the sensitivities follow from linear systems with the converged microscopic tangent. The reference implementation uses one nonlinear equilibrium solution followed by three sensitivity right-hand sides, one per strain component. Independent finite-difference checks perturb the input and re-solve equilibrium. Such checks assess the constitutive linearization, not the physical accuracy of the homogenization assumptions. The same consistency requirement applies to each surrogate: its tangent must differentiate its own stress evaluation.

2.5 Macroscopic equilibrium and the constitutive interface

The pair $(\mathbf{s}, \mathbf{D})$ is the common interface through which every constitutive evaluator enters the structural problem. Let $\mathbf{U}$ collect the free macroscopic nodal displacements, and let $\mathbf{e}_g(\mathbf{U})$ be the engineering Green strain at integration point g . With $\mathbf{B}_g = \partial \mathbf{e}_g / \partial \mathbf{U}$, reference-volume integration weights Ω_g , and dead external loads $\mathbf{f}_{\text{ext}}$, equilibrium reads

$$\mathbf{R}_M(\mathbf{U}) = \sum_g \Omega_g \mathbf{B}_g^T \mathbf{s}_g(\mathbf{e}_g(\mathbf{U})) - \mathbf{f}_{\text{ext}} = \mathbf{0}. \quad (11)$$

Its tangent contains material and geometric contributions,

$$\mathbf{K}_M = \sum_g \Omega_g \left[\mathbf{B}_g^T \mathbf{D}_g \mathbf{B}_g + \sum_{\alpha=1}^3 (s_g)_\alpha \frac{\partial^2 (e_g)_\alpha}{\partial \mathbf{U} \partial \mathbf{U}} \right]. \quad (12)$$

The first term uses the constitutive tangent; the second arises because the Green strain is nonlinear in the displacements. A structural Newton solver therefore needs both $\mathbf{s}_g$ and $\mathbf{D}_g$, irrespective of how the constitutive response is computed.

The nested FOM–FE² evaluation provides the reference instance of this interface: every constitutive query solves the full periodic RVE problem defined above. Section 3 next constructs effective energies that return $\mathbf{s}$ and $\mathbf{D}$ by differentiation and therefore remove the microscopic solve from the online stage. Section 4 then introduces the intrusive reduced references. Its iterative HPROM variants retain microscopic equilibrium in reduced coordinates and replace full assembly by hyperreduced integration, whereas the direct decoder bypasses that equilibrium solve. Despite these different routes, all models supply the same interface while the macroscopic mesh, loading, and equilibrium algorithm remain unchanged.

Holding that structural problem fixed makes the nested FOM–FE² solution the reference for measuring surrogate errors at fixed macro and micro discretizations. It is not a direct numerical simulation resolving every pore in the coupon. Homogenization error, boundary-layer effects, and mesh-discretization error are therefore distinct from the surrogate errors reported below. The cell itself is an ideal periodic material model, without a claim of statistical representativeness for a random microstructure. Accordingly, FE² denotes the nested calculation in which each constitutive query solves an RVE problem; a direct learned constitutive evaluation supplies the same interface but is not itself a microscopic finite element solve.

3 Learnable paired-feature polyconvex energies

The non-intrusive route developed here replaces the repeated online solution of the microscopic problem by a learned effective stored energy. Its central constitutive relation is

$$\mathbf{s}(\mathbf{e}) = \nabla_{\mathbf{e}} W(\mathbf{e}), \quad (13)$$

where $\mathbf{s}$ is the stress work-conjugate to the engineering-strain vector $\mathbf{e}$. Thus all stress components derive from a common potential, and the Hessian of W supplies a major-symmetric material tangent. Fitting stress components independently does not, in general, enforce the cross-derivative identities required by such a potential [12]. This observation motivates the energy-based formulation adopted throughout this section.

One could stop at this point and construct a flexible hyperelastic network without restricting the curvature of its learned potential. We use precisely this option as the *Unconstrained energy* reference. It takes objective strain measures as inputs and obtains stress through Eq. (13). Beyond potential structure and objectivity, the only mechanical conditions imposed by construction are the reference conditions $W(\mathbf{I}) = 0$ and $\partial_{\mathbf{e}} W(\mathbf{0}) = \mathbf{0}$. The first fixes the arbitrary additive energy constant; the second is enforced by an affine correction that cancels the stress at the undeformed state, following the consistency construction of As’ad et al. [12].²

The Unconstrained energy model shows what can be represented before stronger finite-strain restrictions are introduced. Potential structure and reference normalization alone do not control deformation-gradient curvature or prevent finite energy as $J \rightarrow 0^+$. We therefore pursue polyconvexity by constructing the energy as a convex function of independent deformation minors, together with analytic terms that control the reference state and growth. The learned contribution combines a common family of paired directional features with one of two monotone convex cores: an input-convex neural network (ICNN) or an input-convex Kolmogorov–Arnold network (ICKAN). These cores provide alternative representations while preserving the same feature construction and mechanical guarantees.

²No convexity or growth condition is imposed on this reference.

The remainder of this section makes these distinctions precise. We first state the mechanical requirements and separate potential structure from finite-strain stability. We then define the reference model, construct the paired features and convex cores, and establish the properties of the complete energy. Throughout, $\mathbf{F} \in \text{GL}^+(2)$, $J = \det \mathbf{F} > 0$, and derivatives with respect to deformation are taken with the learned parameters fixed.

3.1 Admissibility criteria and scope of the guarantees

A small prediction error is not sufficient to define an admissible hyperelastic surrogate. We therefore separate six requirements relevant to the present construction. The requirements themselves are standard in constitutive PANN formulations [13, 14]; their grouping and labels below are specific to this paper:

- (C1) *Potential structure*: stress is the derivative of a single-valued stored energy in work-conjugate variables.
- (C2) *Objectivity*: a superposed rigid rotation does not change that energy.
- (C3) *Stress symmetry and anisotropy*: the second Piola stress is symmetric, while directional dependence is allowed without imposing an unverified material symmetry group.
- (C4) *Reference normalization*: the undeformed state has zero energy and zero stress.
- (C5) *Polyconvexity*: the energy admits a convex representation in deformation minors, providing a sufficient rank-one convexity condition.
- (C6) *Growth*: energy diverges under volume collapse, unbounded volume expansion, and unbounded distortion at bounded positive volume.

This list is an organizing device, not a universal checklist or a claim that polyconvexity is necessary for every physical material. The Unconstrained energy reference satisfies (C1)–(C4) but carries no construction-level claim of (C5) or (C6). The ICNN and ICKAN models proposed here satisfy all six criteria by construction for every finite parameter set that obeys the stated architectural constraints.³

These statements rest on different levels of evidence. Constitutive identities follow from the representation itself; sufficient bounds hold under their stated assumptions; and sampled numerical checks verify an implementation but do not replace a proof. The discussion below maintains this distinction while examining each criterion.

Potential structure (C1). To ensure mutually compatible stresses, we derive all stress components from a single scalar potential. Equation (13) gives the resulting work-conjugate stress. It also implies that the work around every closed strain path Γ vanishes:

$$\oint_{\Gamma} \mathbf{s} \cdot d\mathbf{e} = 0. \quad (14)$$

For the isothermal purely elastic setting considered here, (C1) likewise gives $\mathbf{s} \cdot \dot{\mathbf{e}} - \dot{W} = 0$: the mechanical dissipation vanishes. Moreover, the stress Jacobian is the Hessian of W and therefore obeys the cross-derivative identities $\partial s_i / \partial e_j = \partial s_j / \partial e_i$. These statements follow from the scalar-potential representation itself, not from a small error at sampled states.

Objectivity and the three meanings of symmetry (C2)–(C3). Objectivity must remove dependence on the spatial observer without inadvertently imposing material isotropy. For a superposed proper rotation $\mathbf{Q}$, it requires

$$W(\mathbf{Q}\mathbf{F}) = W(\mathbf{F}), \quad \mathbf{P}(\mathbf{Q}\mathbf{F}) = \mathbf{Q}\mathbf{P}(\mathbf{F}). \quad (15)$$

Using $\mathbf{C} = \mathbf{F}^T \mathbf{F}$ as an energy input enforces the first identity, because a superposed spatial rotation leaves $\mathbf{C}$ unchanged; differentiating the objective energy then gives the stress transformation. Objectivity does not itself imply either potential structure or material isotropy. For the energy models considered here, the symmetry $\mathbf{S} = \mathbf{S}^T$ follows from differentiation with respect to symmetric strain. It implies a symmetric Cauchy stress $\boldsymbol{\sigma} = J^{-1} \mathbf{F} \mathbf{S} \mathbf{F}^T$, as required by angular-momentum balance in the present continuum.

Objectivity still permits anisotropy. Material symmetry is a different invariance: a right action $\mathbf{F} \mapsto \mathbf{F} \mathbf{Q}_m$ leaves the energy unchanged only when $\mathbf{Q}_m$ belongs to the material symmetry group. Section 3.3 will implement this distinction using directions attached to the reference material; a superposed spatial rotation

³Global nonnegative energy is not part of (C1)–(C6); the saved checkpoints are audited separately in Section 5.5 using the sufficient bound in Appendix C.

acts on $\mathbf{F}$ from the left and does not alter the directional stretch measured in that reference. Finally, the constitutive tangent $\mathbf{D} = \partial \mathbf{s} / \partial \mathbf{e} \in \mathbb{R}^{3 \times 3}$ measures how the three stress components change with engineering strain, with entries $D_{ij} = \partial s_i / \partial e_j$. For an energy-based law it is the Hessian $\partial_{\mathbf{e}\mathbf{e}}^2 W$. Its *major symmetry*, $\mathbf{D} = \mathbf{D}^T$, is a third property: it follows from a twice differentiable scalar energy, not merely from a symmetric stress tensor.

Reference normalization (C4). To make the learned energy consistent with the undeformed reference state, we require both zero energy and zero stress:

$$W(\mathbf{I}) = 0, \quad \mathbf{s}(\mathbf{0}) = \mathbf{0}, \quad (16)$$

where $\mathbf{s}(\mathbf{0}) = \mathbf{0}$ is equivalent to $\mathbf{P}(\mathbf{I}) = \mathbf{0}$ in the undeformed state. The first condition fixes an arbitrary additive energy constant. The second is a substantive stress condition, and the admissible correction depends on the model class. Section 3.2 subtracts the raw reference value and a term linear in $\mathbf{E}$ for the Unconstrained energy reference. That correction preserves its scalar potential, but an arbitrary linear function of $\mathbf{E}$ need not preserve polyconvexity in $\mathbf{F}$. Section 3.5 therefore uses a different correction for the proposed models: it is affine in independent deformation minors and thus preserves convexity in those variables.

Polyconvexity and finite-strain stability (C5). Potential structure alone does not control the energy curvature under changes in deformation. Direct convexity in $\mathbf{F}$ is generally too restrictive for objective finite-strain energies with singular-volume growth. Polyconvexity instead imposes convexity in independent deformation minors [17, 18]. In two dimensions,⁴ this means that the energy admits the representation

$$W(\mathbf{F}) = \mathcal{P}(\mathbf{F}, \det \mathbf{F}), \quad \mathcal{P} \text{ convex in its independent arguments.} \quad (17)$$

Here $\mathcal{P}$ is defined on the convex set $\mathbb{R}^{2 \times 2} \times (0, \infty)$, treating its matrix and scalar arguments as independent; only its evaluation in W sets the second argument equal to $\det \mathbf{F}$. The cofactor is linear in the entries of a 2×2 matrix, so it need not be introduced as an additional independent argument. Convexity of $\mathcal{P}$ implies that W cannot curve downward along an admissible rank-one perturbation of $\mathbf{F}$. Appendix A derives this consequence and states its differential form.

To distinguish the present construction from alternatives that impose convexity in strain space [12, 15], we emphasize that convexity in $\mathbf{C}$ or $\mathbf{E}$ does not imply polyconvexity. The reason is that $\mathbf{E} = \frac{1}{2}(\mathbf{F}^T \mathbf{F} - \mathbf{I})$ depends quadratically on $\mathbf{F}$; Appendix A provides the corresponding derivation and an analytical counterexample.

Growth and the limits of the guarantees (C6). Polyconvexity alone does not prevent volume collapse or unbounded deformation. We therefore require $W \rightarrow \infty$ as $J \rightarrow 0^+$, as $J \rightarrow \infty$, and as $\|\mathbf{F}\| \rightarrow \infty$ while J remains bounded away from zero and infinity. These growth conditions contribute to coercivity but do not alone prove existence of a minimizer: the standard existence argument also requires admissible boundary data and finite-energy competitors [18]. Nor does a constitutive barrier guarantee convergence of an unconstrained Newton iteration. Architectural guarantees and numerical solver behavior are therefore reported separately.

3.2 Unconstrained energy reference

The most flexible energy model used in the comparison is a smooth neural network $\widehat{W}_\theta$ acting on

$$\mathbf{x}(\mathbf{e}) = (C_{11} - 1, C_{22} - 1, C_{12}, J - 1)^T. \quad (18)$$

Here $\mathbf{C} = \mathbf{I} + 2\mathbf{E}$ and $J = \sqrt{\det \mathbf{C}}$, so the four objective inputs are functions of the same three engineering-strain components; they do not introduce an additional state variable. Denote the resulting raw energy by $\widetilde{W}(\mathbf{e}) = \widehat{W}_\theta(\mathbf{x}(\mathbf{e}))$. A generic network does not satisfy the reference conditions automatically, so we define the normalized energy as

$$W_{\text{unc}}(\mathbf{e}) = \widetilde{W}(\mathbf{e}) - \widetilde{W}(\mathbf{0}) - \partial_{\mathbf{e}} \widetilde{W}(\mathbf{0}) \cdot \mathbf{e}. \quad (19)$$

At $\mathbf{e} = \mathbf{0}$, the first two terms cancel and the linear term vanishes, giving $W_{\text{unc}}(\mathbf{0}) = 0$. Differentiation gives $\partial_{\mathbf{e}} W_{\text{unc}}(\mathbf{e}) = \partial_{\mathbf{e}} \widetilde{W}(\mathbf{e}) - \partial_{\mathbf{e}} \widetilde{W}(\mathbf{0})$, and therefore zero reference stress. Because the correction is affine in $\mathbf{e}$, it does not change the Hessian of the raw network. Equation (19) thus fixes the reference energy and stress while leaving the learned curvature unrestricted.

Using components of $\mathbf{C}$ in a material reference basis ensures objectivity under superposed spatial rotations but does not impose isotropy: the network may depend differently on C_{11} , C_{22} , and C_{12} . A change of material reference axes must transform these tensor components and the represented law consistently. This model is used only as the flexible energy-based reference for interpreting the constrained ICNN and ICKAN results.

⁴The standard three-dimensional definition uses $(\mathbf{F}, \text{cof } \mathbf{F}, \det \mathbf{F})$. The two-dimensional guarantee established here does not imply an admissible extension to arbitrary out-of-plane deformation.

3.3 Feature design: expressiveness and admissibility

The first design question is how to describe deformation without erasing the effective anisotropy of the RVE. Classical invariants provide the natural starting point. For a three-dimensional right Cauchy–Green tensor $\mathbf{C}_3$, the familiar first, second, and third principal invariants are

$$I_1^{(3)} = \text{tr } \mathbf{C}_3, \quad I_2^{(3)} = \frac{1}{2}[(\text{tr } \mathbf{C}_3)^2 - \text{tr}(\mathbf{C}_3^2)], \quad I_3^{(3)} = \det \mathbf{C}_3. \quad (20)$$

They describe principal stretches without recording their orientation relative to the material. An energy depending only on these scalars is therefore isotropic. In the present two-dimensional setting, the corresponding two independent principal invariants are $I_1 = \text{tr } \mathbf{C}$ and $\det \mathbf{C} = J^2$; a third independent principal invariant is not available.

To retain directional information, one can supplement the isotropic measures with a structural direction $\mathbf{d}$. The familiar anisotropic stretch invariant is

$$I_4(\mathbf{d}) = \mathbf{d}^T \mathbf{C} \mathbf{d} = \mathbf{C} : (\mathbf{d} \otimes \mathbf{d}) = \|\mathbf{F} \mathbf{d}\|^2, \quad \|\mathbf{d}\| = 1. \quad (21)$$

It measures squared stretch along a material direction and remains unchanged by a superposed spatial rotation. Such directional invariants are standard ingredients of anisotropic hyperelastic energies [40, 41]. Here $\mathbf{d}$ need not represent a physical fiber: learned directions provide objective scalar descriptors of the RVE’s effective anisotropy. In two dimensions, $\mathbf{d}$ is naturally accompanied by its perpendicular $\mathbf{d}^\perp$. For this orthonormal pair,

$$I_4(\mathbf{d}) + I_4(\mathbf{d}^\perp) = I_1, \quad \|\text{cof } \mathbf{F} \mathbf{d}\|^2 = I_4(\mathbf{d}^\perp). \quad (22)$$

The first identity shows that the pair resolves how I_1 is distributed between two perpendicular directions: their sum recovers I_1 , while their individual values retain directional information. The second shows that a cofactor-based directional measure reads the stretch in the perpendicular direction rather than introducing an independent third in-plane descriptor. These observations motivate the paired directional building block used below.

Objective directional descriptors solve only the first part of the design problem. Because the downstream core is constrained to be nondecreasing in every input, the feature coordinates also determine which deformation states the core must order. The feature map must therefore be designed together with that monotonicity constraint. Consider a feature vector $\mathbf{z}(\mathbf{F})$ and a scalar core H nondecreasing in each input. For any two deformation states,

$$\mathbf{z}(\mathbf{F}_B) \geq \mathbf{z}(\mathbf{F}_A) \text{ componentwise} \implies H(\mathbf{z}(\mathbf{F}_B)) \geq H(\mathbf{z}(\mathbf{F}_A)). \quad (23)$$

This implication holds regardless of the network’s width. If the complete constitutive energy is $W = H(\mathbf{z}) + W_{\text{corr}}$, where W_{corr} collects the analytic normalization and growth terms, the ordering constrains the neural contribution $W - W_{\text{corr}}$ rather than the complete energy. The feature map and analytic terms must therefore be considered together.

This observation explains why the choice of an admissible feature construction can matter more than the capacity of its convex core. Changing the number of features does not change the number of independent strains: the features remain functions of the same three engineering-strain components. They embed that state into variables on which the network is constrained to be monotone and convex. Additional features can enlarge the admissible response family, but their number alone does not prove completeness.

We therefore seek features that retain directional stretch, reshape the componentwise ordering seen by the monotone core, and remain convex in the independent variables $(\mathbf{F}, J)$. The paired construction below achieves these three objectives by coupling directional stretch to volume through admissible determinant exponents. A feature can consequently decrease along a path when its denominator grows sufficiently, even if a numerator increases. Its reference derivative is also controlled analytically, which will allow the reference stress to be removed without sacrificing convexity.

3.3.1 Paired directional features

Let m denote the number of scalar features. For each $i = 1, \dots, m$, let $\mathbf{d}_i = (\cos \theta_i, \sin \theta_i)^T$ and $\mathbf{d}_i^\perp = (-\sin \theta_i, \cos \theta_i)^T$. Define

$$\begin{aligned} z_i(\mathbf{F}, J) &= \frac{I_4(\mathbf{d}_i)^{p_i}}{p_i J^{b_i}} + \frac{I_4(\mathbf{d}_i^\perp)^{q_i}}{q_i J^{c_i}} \\ &= \frac{\|\mathbf{F} \mathbf{d}_i\|^{2p_i}}{p_i J^{b_i}} + \frac{\|\mathbf{F} \mathbf{d}_i^\perp\|^{2q_i}}{q_i J^{c_i}}, \end{aligned} \quad (24)$$

where θ_i sets the orientation of the reference directions, p_i, q_i are the powers of the squared-stretch invariants, and b_i, c_i control each contribution's coupling to the area ratio J . These roles provide expressiveness, but the exponents cannot be chosen arbitrarily if the feature is to remain convex. They satisfy

$$p_i, q_i \geq \frac{1}{2}, \quad 0 \leq b_i \leq 2p_i - 1, \quad 0 \leq c_i \leq 2q_i - 1. \quad (25)$$

The lower bounds on p_i and q_i make the corresponding powers of the directional norms convex, while the upper bounds on b_i and c_i prevent their coupling with J from introducing a negative-curvature direction. Appendix B derives these conditions from the Hessian of a single stretch-determinant contribution. In the learned-feature variants, θ_i and the four exponent parameters are optimized, with p_i, q_i, b_i, c_i subject to Eq. (25); the fixed-feature controls of Section 5.2 hold the same initialized feature parameters constant while training the core. We use $\text{softplus}(x) = \log(1 + e^x)$ and $\text{sigmoid}(x) = (1 + e^{-x})^{-1}$, whose ranges for $x \in \mathbb{R}$ are $(0, \infty)$ and $(0, 1)$, respectively. With unconstrained parameters $\hat{p}_i, \hat{q}_i, \hat{b}_i, \hat{c}_i \in \mathbb{R}$, set

$$\begin{aligned} p_i &= \frac{1}{2} + \text{softplus}(\hat{p}_i), & b_i &= (2p_i - 1) \text{sigmoid}(\hat{b}_i), \\ q_i &= \frac{1}{2} + \text{softplus}(\hat{q}_i), & c_i &= (2q_i - 1) \text{sigmoid}(\hat{c}_i). \end{aligned} \quad (26)$$

The angles require no admissibility bound: for any fixed θ_i , the maps $\mathbf{F} \mapsto \mathbf{F}\mathbf{d}_i$ and $\mathbf{F} \mapsto \mathbf{F}\mathbf{d}_i^\perp$ remain linear in $\mathbf{F}$. The transformations in Eq. (26) preserve the exponent inequalities during training, so a penalty loss is not needed to enforce them. For finite real values of the unconstrained parameters, the bounds in Eq. (25) hold strictly; boundary values are approached as limits. Each z_i is one scalar feature containing two perpendicular contributions. The two contributions need not have equal exponents. Consequently, pairing perpendicular directions does not impose invariance under their interchange.

The construction was also designed to simplify reference normalization. Its paired and scaled contributions make the reference strain gradient of every feature isotropic. Consequently, the neural contribution has an isotropic reference stress, which a single scalar determinant correction can cancel in Section 3.5. The following result gathers the objectivity, convexity, and reference-state properties needed there.

Proposition 1. *Under the constraints in Eq. (25), each feature in Eq. (24) is objective and convex in the independent variables $(\mathbf{F}, J)$ on $J > 0$. Its value and engineering-strain gradient at the identity, evaluated with $J = \det \mathbf{F}$, are*

$$z_i^0 = \frac{1}{p_i} + \frac{1}{q_i}, \quad \partial_{\mathbf{e}} z_i|_{\mathbf{0}} = \rho_i(1, 1, 0)^T, \quad \rho_i = 2 - \frac{b_i}{p_i} - \frac{c_i}{q_i}. \quad (27)$$

The proof is given in Appendix B. The mechanism behind the reference gradient is that $1/p_i$ and $1/q_i$ cancel the stretch-power multipliers upon differentiation. The two perpendicular directional contributions then sum to $2\mathbf{I}$, while the determinant terms contribute $-(b_i/p_i + c_i/q_i)\mathbf{I}$. Thus ρ_i is the scalar amplitude of an isotropic reference gradient: its two normal components are equal and its shear component is zero. This reference property does not imply isotropy of the response away from the reference state.

The raw features are now mechanically admissible, but their numerical magnitudes may differ. Before passing them to the core, we therefore use $y_i = (z_i - z_i^{\text{cen}})/s_i$, where z_i^{cen} is a centering value and $s_i > 0$ is a scaling factor. Both are independent of deformation for any fixed parameter set. This positive affine conditioning preserves both convexity in independent $(\mathbf{F}, J)$ and the componentwise ordering required by the monotone core. The choices used in the constitutive assessments are specified in Section 5.2.

This feature design has a precise but limited purpose: it supplies objective, expressive, and admissible inputs to the constrained core. Learning the directions does not amount to identifying or enforcing a complete material symmetry group, and no completeness or universal-approximation theorem is asserted for a finite collection of these features followed by a monotone convex core.

3.4 Monotone ICNN and integrated-hat ICKAN cores

The feature construction of Section 3.3 resolves the input side of the constitutive design: each z_i is objective and convex in the independent variables $(\mathbf{F}, J)$. The remaining question is how a learned core may combine these inputs without undoing that convexity. The core receives the conditioned vector $\mathbf{y} = (y_1, \dots, y_m)^T$, with $y_i = (z_i - z_i^{\text{cen}})/s_i$. Because $s_i > 0$, this affine conditioning preserves the convexity and componentwise ordering of the physical features.

A convex core alone is not sufficient. Composition with the convex feature map preserves convexity only when the core is also componentwise nondecreasing [13, 14]. We enforce these same two properties in two different ways. The ICNN learns nonnegative linear maps and applies fixed softplus activations at its hidden nodes. The ICKAN instead learns a convex, nondecreasing scalar function on every edge and sums the incoming contributions at each node. The following definitions make these two mechanisms precise; Figure 2 then compares their multilayer architectures.




Figure 2: Multilayer monotone cores mapping the conditioned paired features $\mathbf{y}$ to the scalar contribution $\hat{H}(\mathbf{y})$. (a) The ICNN learns nonnegative linear maps and applies fixed softplus activations at hidden nodes. (b) The ICKAN learns one convex, nondecreasing scalar function per edge and sums them at nodes. Purple dashed paths denote nonnegative maps of the complete input; ellipses denote intervening layers. Glyph counts indicate where each operation acts, not deployed widths or parameter counts. Unrestricted biases are omitted, and the diagram represents the neural cores rather than the complete energy.

3.4.1 Monotone ICNN

Let L denote the number of hidden layers and $\mathbf{a}^{(\ell)}$ the vector of activations in layer ℓ , for $\ell = 1, \dots, L$. We write $\hat{H}(\mathbf{y})$ for the scalar neural core acting on the conditioned features $\mathbf{y}$; the complete energy additionally includes the reference-normalization and growth terms of Section 3.5. The deployed ICNN has the form

$$\begin{aligned}
 \mathbf{a}^{(1)} &= \phi(\mathbf{U}_1 \mathbf{y} + \mathbf{b}_1), \\
 \mathbf{a}^{(\ell)} &= \phi(\mathbf{W}_\ell \mathbf{a}^{(\ell-1)} + \mathbf{U}_\ell \mathbf{y} + \mathbf{b}_\ell), \quad \ell = 2, \dots, L, \\
 \hat{H}(\mathbf{y}) &= \mathbf{v}^T \mathbf{a}^{(L)} + \mathbf{w}^T \mathbf{y} + b_{\text{out}},
 \end{aligned} \tag{28}$$

where $\phi = \text{softplus}$ is applied componentwise. Its derivatives satisfy $\phi'(x) = \text{sigmoid}(x) > 0$ and $\phi''(x) = \text{sigmoid}(x)[1 - \text{sigmoid}(x)] > 0$. Thus the activation is smooth, increasing, and convex; its smoothness also makes the core C^2 , as required for a continuous constitutive tangent. The matrices $\mathbf{U}_\ell$ and $\mathbf{W}_\ell$ contain the weights from $\mathbf{y}$ and the preceding hidden layer, respectively, and $\mathbf{b}_\ell$ is the layer bias vector. The output weight vectors $\mathbf{v}$ and $\mathbf{w}$ act on $\mathbf{a}^{(L)}$ and $\mathbf{y}$, respectively; b_{out} is an additive scalar output bias.

Every entry of $\mathbf{W}_\ell, \mathbf{U}_\ell, \mathbf{v}, \mathbf{w}$ is nonnegative, enforced by softplus transformations of unconstrained weight parameters. These signs play two related but distinct roles. The matrices $\mathbf{W}_\ell$ and the output vector $\mathbf{v}$ multiply convex hidden activations, so their nonnegativity preserves convexity as well as monotonicity. The maps $\mathbf{U}_\ell \mathbf{y}$ and $\mathbf{w}^T \mathbf{y}$ are affine and therefore do not affect convexity for either sign, but nonnegative $\mathbf{U}_\ell$ and $\mathbf{w}$ are required for componentwise monotonicity. Biases are unrestricted because adding constants to a preactivation preserves both properties. Hidden biases can change the core's derivatives with respect to $\mathbf{y}$, whereas b_{out} only shifts $\hat{H}$ and leaves its derivatives unchanged.

The preservation argument proceeds layer by layer. The first hidden layer is convex because softplus is convex and its argument is affine, and it is nondecreasing because $\mathbf{U}_1$ is nonnegative. Assuming that $\mathbf{a}^{(\ell-1)}$ is convex and nondecreasing, the nonnegative map $\mathbf{W}_\ell \mathbf{a}^{(\ell-1)}$, the nonnegative affine input map, and composition with ϕ preserve both properties in layer ℓ . The same argument applies to the output, proving by induction that $\hat{H}$ is convex and componentwise nondecreasing. This is the ICNN composition argument of Amos et al. [42], supplemented here by nonnegative direct-input maps to guarantee monotonicity.

3.4.2 Monotone integrated-hat ICKAN

Unlike the ICNN, which applies a fixed activation after an affine weighted sum at each hidden node, the ICKAN assigns a learned scalar function to each connection (edge) between layers. Each node sums the transformed contributions from its incoming connections. To make the relation between these edge functions and the scalar core explicit, set $\mathbf{h}^{(0)} = \mathbf{y}$ and let n_ℓ be the width of layer ℓ , with $n_0 = m$. The hidden layers and scalar output are

$$\begin{aligned} h_j^{(\ell)} &= \sum_{i=1}^{n_{\ell-1}} f_{ji}^{(\ell)}(h_i^{(\ell-1)}) + b_j^{(\ell)}, & j = 1, \dots, n_\ell, \quad \ell = 1, \dots, L, \\ \hat{H}(\mathbf{y}) &= \sum_{i=1}^{n_L} f_i^{\text{out}}(h_i^{(L)}) + \mathbf{w}^T \mathbf{y} + b_{\text{out}}, & \mathbf{w} \geq \mathbf{0}. \end{aligned} \quad (29)$$

Thus $f_{ji}^{(\ell)}$ transforms one scalar carried by an edge, $h_j^{(\ell)}$ is the sum formed at a node, and $\hat{H}(\mathbf{y})$ is the complete ICKAN core. The nonnegative linear skip $\mathbf{w}^T \mathbf{y}$ connects the conditioned input directly to the scalar output. It remains to specify the functions appearing on the edges. For any $f_{ji}^{(\ell)}$ or f_i^{out} in Eq. (29), let x denote its scalar input. Each edge uses the common form

$$f(x) = ax + \sum_k c_k B_k(x), \quad a, c_k \geq 0, \quad (30)$$

where a is the learned linear slope, c_k are learned weights of the spline bases B_k , and k indexes these bases. Softplus transformations of unconstrained parameters impose the signs of a and c_k . Additive constants are represented once, through the unrestricted node biases in Eq. (29), rather than redundantly on every incoming edge.

The remaining design task is to learn the shape of each edge without allowing negative slope or curvature. Rather than check these inequalities at sampled inputs after training, we parameterize the edge curvature directly. Each B_k'' is a nonnegative triangular hat centered at a knot κ_k . The knots form a fixed uniform grid with spacing $h > 0$; neither their locations nor the spacing are optimized. Writing $s = (x - \kappa_k)/h$, one curvature basis is

$$B_k''(x) = \max(1 - |s|, 0). \quad (31)$$

Integrating twice from the left conditions $B_k(x) = B_k'(x) = 0$ for $x \leq \kappa_k - h$ gives

$$B_k(x) = h^2 \begin{cases} 0, & s \leq -1, \\ (s+1)^3/6, & -1 < s \leq 0, \\ (-s^3 + 3s^2 + 3s + 1)/6, & 0 < s < 1, \\ s, & s \geq 1. \end{cases} \quad (32)$$

Figure 3 shows both levels of the construction: integration of one curvature basis and assembly of a learned edge from several bases.

The value and first two derivatives in Eq. (32) match at every join, and B_k'' vanishes at the outer joins. Thus B_k is globally C^2 . Its curvature is localized to $[\kappa_k - h, \kappa_k + h]$, but the integrated basis is not: it has a zero left tail and a linear right tail. Moreover, $B_k'(-\infty) = 0$ and $B_k'' \geq 0$ imply $B_k' \geq 0$. The complete edge therefore satisfies

$$f''(x) = \sum_k c_k B_k''(x) \geq 0, \quad f'(x) = a + \sum_k c_k B_k'(x) \geq 0, \quad (33)$$

for every learned parameter set with $a, c_k \geq 0$. Hence each edge is convex and nondecreasing by construction, rather than by a sampled numerical test.

Equation (29) sums these scalar edge functions at each node. Because every edge is convex and nondecreasing, induction across the layers shows that every $h_j^{(\ell)}$, and finally $\hat{H}$, is convex and componentwise nondecreasing in $\mathbf{y}$. The nonnegative input-to-output skip preserves both properties. The fixed knot grid and finite number of bases limit where an individual edge can learn curvature: outside their combined support, the edge has linear tails. These tails are linear in the scalar input x , not in $\mathbf{F}$; the complete energy still contains nonlinear features and analytic growth terms. The grid used in the constitutive assessment is reported in Section 5.2. This core is inspired by the monotonic input-convex KAN framework [48], with the integrated-hat bases and analytic tails specified here.




Figure 3: Integrated-hat parameterization of an ICKAN edge. (a) A localized nonnegative curvature hat integrates to a nonnegative slope and a convex C^2 basis with a zero left tail and a linear right tail. (b) Nonnegative learned coefficients combine several fixed bases into a nonnegative edge curvature; integration gives a nondecreasing slope and a convex, nondecreasing connection. The curves illustrate the analytic construction; they do not represent the deployed grid or parameter count.

3.5 The complete energy and its admissibility guarantees

Sections 3.3 and 3.4 constructed the physical feature vector $\mathbf{z}$, its conditioned representation $\mathbf{y}(\mathbf{z})$, and the scalar neural core $\hat{H}(\mathbf{y})$. We write

$$H(\mathbf{z}) = \hat{H}(\mathbf{y}(\mathbf{z})) \quad (34)$$

for the resulting learned contribution. This term represents the anisotropic response, but it does not by itself set the energy and stress references or prescribe growth at the boundary of the admissible deformation domain.

To complete the construction, we first quantify the residual stress produced by H at the undeformed reference state. Let $\mathbf{z}^0 = \mathbf{z}(\mathbf{I}, 1)$. With the feature-gradient factors ρ_i from Eq. (27), define the core sensitivities and the scalar amplitude of this residual stress as

$$h_i = \partial_{z_i} H(\mathbf{z}^0) \geq 0, \quad r = \sum_i h_i \rho_i. \quad (35)$$

For a fixed parameter set, $H(\mathbf{z}^0)$ and r are constants with respect to deformation. The complete stored energy then combines the learned contribution with analytic reference-normalization and growth terms:

$$\begin{aligned} W(\mathbf{F}) = & \underbrace{H(\mathbf{z}) - H(\mathbf{z}^0)}_{\text{(i) } W_{\text{learn}}} + \underbrace{[-r(J-1)]}_{\text{(ii) } W_{\text{corr}}} \\ & + \underbrace{\alpha(J-1 - \log J) + \frac{\beta}{2}(J-1)^2}_{\text{(iii) } W_{\text{vol}}} \\ & + \underbrace{\varepsilon \left(\frac{\text{tr } \mathbf{C}}{2} - 1 - \log J \right)}_{\text{(iv) } W_{\text{dist}}}, \quad \alpha, \beta, \varepsilon > 0. \end{aligned} \quad (36)$$

Here $J = \det \mathbf{F}$, $\mathbf{C} = \mathbf{F}^T \mathbf{F}$, and $\mathbf{z} = \mathbf{z}(\mathbf{F}, J)$. The positive energy-density coefficients α , β , and ε weight the analytic growth terms. For either core, the composition in Eq. (34) is twice continuously differentiable, convex, and componentwise nondecreasing in $\mathbf{z}$. The underbraces name the four contributions used below: W_{learn} carries the learned anisotropic response, while W_{corr} , W_{vol} , and W_{dist} impose the reference and growth behavior.

Each contribution solves a distinct constitutive problem:

- (i) **Learned response and energy reference (W_{learn}).** At $\mathbf{F} = \mathbf{I}$, the features equal $\mathbf{z}^0$, and hence

$$W(\mathbf{I}) = H(\mathbf{z}^0) - H(\mathbf{z}^0) = 0. \quad (37)$$

The remaining terms also vanish because $J = 1$ and $\text{tr } \mathbf{C} = 2$. Since $H(\mathbf{z}^0)$ is constant with respect to deformation, its subtraction fixes the energy reference without changing the stress or tangent.

- (ii) **Reference-stress correction** (W_{corr}). A constant subtraction cannot change stress. By the chain rule and Eq. (27), the learned contribution instead has the reference value

$$\nabla_{\mathbf{e}} H(\mathbf{z})|_0 = \sum_i h_i \nabla_{\mathbf{e}} z_i|_0 = \sum_i h_i \rho_i(1, 1, 0)^T = r(1, 1, 0)^T. \quad (38)$$

The identity $J = \sqrt{(1 + 2E_{11})(1 + 2E_{22}) - (2E_{12})^2}$ gives $\nabla_{\mathbf{e}} J|_0 = (1, 1, 0)^T$. Holding r fixed when differentiating therefore yields

$$\nabla_{\mathbf{e}} [H(\mathbf{z}) - r(J - 1)]|_0 = r(1, 1, 0)^T - r(1, 1, 0)^T = \mathbf{0}. \quad (39)$$

A single scalar r suffices because the paired features make only the *reference* stress isotropic; the learned response away from $\mathbf{F} = \mathbf{I}$ remains anisotropic. The correction is zero in value at $J = 1$ and, being affine in the independent determinant variable, preserves convexity for either sign of r . This affine determinant normalization follows the strategy of Linden et al. [14]; Eq. (27) makes it applicable to the present directional feature family.

- (iii) **Volumetric growth** (W_{vol}). Its two parts control opposite limits: $\alpha(J - 1 - \log J)$ diverges as $J \rightarrow 0^+$, whereas $\beta(J - 1)^2/2$ grows quadratically as $J \rightarrow \infty$. At the reference,

$$W_{\text{vol}}(1) = 0, \quad W'_{\text{vol}}(J) = \alpha(1 - J^{-1}) + \beta(J - 1), \quad W'_{\text{vol}}(1) = 0. \quad (40)$$

Thus this term changes neither the reference energy nor the reference stress. Both parts are needed: the quadratic term alone would remain finite at collapse, with limit $\beta/2$.

- (iv) **Distortional growth at bounded area** (W_{dist}). A function of J alone cannot detect a change of shape that preserves area. For example, with $\lambda > 0$,

$$\mathbf{F} = \text{diag}(\lambda, \lambda^{-1}), \quad J = 1, \quad \varepsilon \left(\frac{\text{tr } \mathbf{C}}{2} - 1 - \log J \right) = \frac{\varepsilon}{2} (\lambda^2 + \lambda^{-2} - 2). \quad (41)$$

This contribution diverges as $\lambda \rightarrow \infty$, although the area remains constant. More generally, $\text{tr } \mathbf{C} = \|\mathbf{F}\|_F^2$ supplies growth under unbounded deformation at bounded positive J . At the reference, the constant -1 sets its value to zero. Using $\mathbf{C} = \mathbf{I} + 2\mathbf{E}$ and $\partial_{\mathbf{E}} J|_0 = \mathbf{I}$, its stress contribution is

$$\partial_{\mathbf{E}} \left[\varepsilon \left(\frac{\text{tr } \mathbf{C}}{2} - 1 - \log J \right) \right] \Big|_0 = \varepsilon [\mathbf{I} - J^{-1} \partial_{\mathbf{E}} J|_0] = \mathbf{0}. \quad (42)$$

Together, W_{learn} and W_{corr} set the learned response and its reference state, while W_{vol} and W_{dist} prescribe growth outside the sampled regime. Zero value and zero first derivative at the reference do not imply zero second derivative: the analytic terms still contribute to the reference stiffness because J is nonlinear in $\mathbf{e}$. The following proposition gathers the guarantees of their complete sum.

Proposition 2. *Let the features satisfy Eq. (25), let the core be either of the C^2 convex, componentwise nondecreasing constructions in Section 3.4, and let $\alpha, \beta, \varepsilon > 0$. Then Eq. (36) defines a C^2 , objective, two-dimensional polyconvex energy on $\text{GL}^+(2)$. Its associated second Piola stress $\mathbf{S}$ is symmetric, its work-conjugate representation is $\mathbf{s} = \nabla_{\mathbf{e}} W$, $W(\mathbf{I}) = 0$, $\mathbf{s}(\mathbf{0}) = \mathbf{0}$, and W diverges under collapse, unbounded expansion, and unbounded distortion at bounded positive area. Thus the construction satisfies (C1)–(C6).*

The proof is given in Appendix B.3. In brief, the convex, componentwise nondecreasing core preserves feature convexity; the remaining terms are constant, affine, or convex in independent $(\mathbf{F}, J)$. The reference identities and growth limits then complete the argument. These guarantees hold for every admissible parameter set and therefore do not depend on the training samples or fitting error. Predictive accuracy and structural deployment are assessed separately in Sections 5.3 and 6. Global nonnegative energy is not an architectural guarantee; it is audited for the saved checkpoints in Section 5.5 using the sufficient bound derived in Appendix C.

The three constitutive constructions carried into the numerical assessment are summarized in Table 1. The ICNN and ICKAN columns include both fixed- and learned-feature variants; learning the feature parameters changes the representation, not the guarantees established above.

Table 1: Constitutive constructions carried into the numerical assessment. Criteria (C1)–(C6) are defined in Section 3.1; guarantees refer to the specific constructions in this section.

	Unconstrained energy	ICNN	ICKAN
Role	Flexible hyperelastic reference	Guaranteed paired-feature energy	Guaranteed paired-feature energy
Inputs	Objective measures from $(\mathbf{C}, J)$	Paired features (fixed or learned)	Paired features (fixed or learned)
Learned scalar map	Smooth energy network	Nonnegative weights; fixed activations	Convex, nondecreasing edge functions
(C1)–(C4)	By construction	By construction	By construction
(C5)–(C6)	Not imposed	By construction	By construction

Section 4 next defines the intrusive microscopic references used in the structural accuracy–cost comparison. Section 5 then evaluates whether the constitutive implementations reproduce the constructions established above and how accurately the resulting energies approximate the RVE response.

4 Projection-based reduced-order models

The learned energies of Section 3 remove the microscopic solve from online constitutive evaluation. To place their structural accuracy and cost against physics-based alternatives, Section 6 also uses three projection-based microscopic surrogates as intrusive references. These models instead approximate the microscopic displacement field while retaining the underlying constituent law. Reduced coordinates decrease the number of equilibrium unknowns; hyperreduction decreases the required element evaluations. We consider an affine HPROM, an iterative HPROM–ANN with adaptive integration, and its direct counterpart D-HPROM–ANN.

Figure 4 provides a visual roadmap. Its first three stages follow the offline construction from RVE snapshots to the affine and nonlinear state representations; the final two compare the three online routes and their common stress–tangent output. The details are introduced in the remainder of this section.

4.1 Constrained microscopic reference

All three reduced evaluators derive from the periodic RVE problem of Section 2. Following the nomenclature of Ref. [24], we refer to its N -dimensional constrained discretization as the high-dimensional model (HDM), or full-order model (FOM). A PROM solves the same equilibrium problem in reduced coordinates, and an HPROM additionally hyperreduces its projected operators.

To express this common equilibrium problem in the coordinates used by the reduced models, we distinguish the physical nodal displacement vector $\mathbf{d} \in \mathbb{R}^{N_{\text{dof}}}$ from the independent state $\mathbf{u} \in \mathbb{R}^N$ obtained after enforcing periodicity and removing rigid translation. In the deployed convention,

$$\mathbf{d}(\mathbf{u}, \mathbf{e}) = \mathbf{T}\mathbf{u} + \mathbf{g}_{\text{per}}(\mathbf{e}), \quad \mathbf{R}(\mathbf{u}; \mathbf{e}) = \mathbf{T}^T \mathbf{f}_{\text{int}}(\mathbf{d}(\mathbf{u}, \mathbf{e})) = \mathbf{0}. \quad (43)$$

Here $\mathbf{T}$ identifies periodic degrees of freedom, and the lifting $\mathbf{g}_{\text{per}}$ supplies the strain-dependent displacement jumps and anchor values defined in Section 2. Thus any independent state $\mathbf{u}$ produces a physical field satisfying the prescribed kinematic constraints. If $\mathbf{K}_{\text{phys}} = \partial_{\mathbf{d}} \mathbf{f}_{\text{int}}$ is the assembled physical stiffness, the constrained HDM tangent is $\mathbf{K} = \partial_{\mathbf{u}} \mathbf{R} = \mathbf{T}^T \mathbf{K}_{\text{phys}} \mathbf{T}$. We use the internal-force sign convention for all residuals.

The reduced bases below are built from snapshots of this affine-retaining state $\mathbf{u}$, rather than from the fluctuation coordinates $\mathbf{a}$ of Eq. (7). If $\mathbf{u}_{\text{aff}}(\mathbf{e})$ denotes the affine displacement at the independent nodes, the two descriptions satisfy $\mathbf{a} = \mathbf{u} - \mathbf{u}_{\text{aff}}(\mathbf{e})$. This coordinate choice leaves the periodic boundary conditions unchanged and preserves load-correlated information used by the strain-informed identification below. Consequently, the reported POD dimension characterizes this representation and its truncation tolerance; it is not an intrinsic material dimension.

4.2 POD approximation and the conventional PROM

The first reduction targets the number of equilibrium unknowns. We collect HDM solutions $\mathbf{u}^s$ at sampled macroscopic strains $\mathbf{e}^s$, with $s = 1, \dots, N_s$, and center them about a reference state $\mathbf{u}_{\text{ref}}$. Their singular




Figure 4: Intrusive reduced references used in the structural comparison. HDM snapshots provide either an affine state representation or a primary–secondary partition in which the ANN predicts the secondary coordinates. Online, HPROM and HPROM–ANN solve projected equilibrium with fixed and adaptive cubature, respectively, whereas D-HPROM–ANN evaluates the nonlinear decoder directly. All three retain microscopic constitutive evaluations and return the same stress–tangent interface.

value decomposition (SVD) is

$$\mathbf{S}_{\text{snap}} = [\mathbf{u}^1 - \mathbf{u}_{\text{ref}}, \dots, \mathbf{u}^{N_s} - \mathbf{u}_{\text{ref}}] = \mathbf{U}_S \mathbf{\Sigma}_S \mathbf{Y}_S^T. \quad (44)$$

Proper orthogonal decomposition (POD) retains the columns of $\mathbf{U}_S$ associated with the largest singular values. Collecting n_{tra} of these modes in the orthonormal reduced-order basis (ROB) $\mathbf{V}_{\text{tra}}$ gives the conventional affine approximation

$$\tilde{\mathbf{u}} = \mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tra}} \mathbf{q}_{\text{tra}}, \quad \mathbf{q}_{\text{tra}} \in \mathbb{R}^{n_{\text{tra}}}, \quad \mathbf{V}_{\text{tra}}^T \mathbf{V}_{\text{tra}} = \mathbf{I}. \quad (45)$$

The PROM therefore replaces the N independent displacement coordinates by $n_{\text{tra}} \ll N$ modal amplitudes. Here *linear PROM* describes only the affine approximation space; the constitutive law and the reduced equilibrium equations remain nonlinear.

The retained dimension is selected from the relative squared reconstruction error. For a Euclidean POD,

$$\epsilon_{\text{POD}}^2(n_{\text{tra}}) = \frac{\sum_{j > n_{\text{tra}}} \sigma_j^2}{\sum_j \sigma_j^2}. \quad (46)$$

This criterion measures how accurately the supplied displacement snapshots are compressed; it is not an error bound for stresses at unsampled strains. Force evaluation, constitutive differentiation, and equilibrium conditioning can make small displacement components important.

Once the approximation space is fixed, equilibrium is enforced by Galerkin projection. At a prescribed strain $\mathbf{e}$, the admissible variations are $\delta \tilde{\mathbf{u}} = \mathbf{V}_{\text{tra}} \delta \mathbf{q}_{\text{tra}}$; requiring zero residual work against them gives the reduced equations and tangent

$$\mathbf{r}_{\text{tra}}(\mathbf{q}_{\text{tra}}; \mathbf{e}) = \mathbf{V}_{\text{tra}}^T \mathbf{R}(\mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tra}} \mathbf{q}_{\text{tra}}; \mathbf{e}) = \mathbf{0}, \quad \mathbf{K}_{\text{tra}} = \mathbf{V}_{\text{tra}}^T \mathbf{K} \mathbf{V}_{\text{tra}}. \quad (47)$$

POD has reduced the nonlinear system to n_{tra} unknowns, but evaluating $\mathbf{R}$ and $\mathbf{K}$ still visits every microscopic element. Hyperreduction, introduced below, addresses this remaining cost.

4.3 Nonlinear reduction through latent-space closure

The POD dimension required to reconstruct a displacement field can exceed the number of coordinates needed to parametrize the loading family. Latent-space closure exploits this distinction: it retains a rich linear span, but learns dependencies among its modal amplitudes so that only a smaller subset remains unknown in equilibrium [22, 24]. We partition an $n_{\text{tot}} = n + \bar{n}$ dimensional ROB into orthonormal primary and secondary blocks,

$$\mathbf{V}_{\text{tot}} = [\mathbf{V}, \bar{\mathbf{V}}], \quad \mathbf{V}^T \mathbf{V} = \mathbf{I}, \quad \bar{\mathbf{V}}^T \bar{\mathbf{V}} = \mathbf{I}, \quad \mathbf{V}^T \bar{\mathbf{V}} = \mathbf{0}, \quad (48)$$

and project each snapshot onto the corresponding coordinates,

$$\mathbf{q}^s = \mathbf{V}^T (\mathbf{u}^s - \mathbf{u}_{\text{ref}}), \quad \bar{\mathbf{q}}^s = \bar{\mathbf{V}}^T (\mathbf{u}^s - \mathbf{u}_{\text{ref}}). \quad (49)$$

The primary coordinates $\mathbf{q} \in \mathbb{R}^n$ remain independent unknowns, whereas an ANN closure $\mathcal{N} : \mathbb{R}^n \rightarrow \mathbb{R}^{\bar{n}}$ predicts the secondary coordinates. The resulting nonlinear decoder and its tangent are

$$\tilde{\mathbf{u}}(\mathbf{q}) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \bar{\mathbf{V}} \mathcal{N}(\mathbf{q}; \boldsymbol{\eta}), \quad \mathbf{B}(\mathbf{q}) = \partial_{\mathbf{q}} \tilde{\mathbf{u}} = \mathbf{V} + \bar{\mathbf{V}} \mathcal{N}_{,\mathbf{q}}. \quad (50)$$

Thus the ANN predicts neither stress nor energy: it predicts the amplitudes needed to reconstruct the microscopic displacement. The secondary modes remain in that field, but no longer appear as independent equilibrium unknowns.

The closure is trained on snapshot pairs $(\mathbf{q}^s, \bar{\mathbf{q}}^s)$ by minimizing

$$\mathcal{L}_{\text{cl}}(\boldsymbol{\eta}) = \frac{1}{|\mathcal{I}_{\text{fit}}| \bar{n}} \sum_{s \in \mathcal{I}_{\text{fit}}} \|\mathcal{N}(\mathbf{q}^s; \boldsymbol{\eta}) - \bar{\mathbf{q}}^s\|_2^2. \quad (51)$$

Input standardization is absorbed into the closure map, while the secondary outputs remain unscaled. Smooth tanh activations provide the decoder derivatives required by projection. This coordinate loss measures displacement reconstruction within the secondary space; it is not a stress- or energy-error guarantee.

The leading POD coefficients need not provide a well-conditioned primary-to-secondary map. We therefore use the input-informed construction of Ref. [39] to rotate the retained POD span and obtain three primary coordinates $\boldsymbol{\xi}$ approximately aligned with the three macroscopic strain components. If $\mathbf{A}_m$ denotes the resulting nonsingular coordinate transformation, the deployed decoder is

$$\tilde{\mathbf{u}}(\boldsymbol{\xi}) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{A}_m \boldsymbol{\xi} + \bar{\mathbf{V}} \hat{\mathcal{N}}(\boldsymbol{\xi}), \quad \mathbf{B}_{\boldsymbol{\xi}} = \partial_{\boldsymbol{\xi}} \tilde{\mathbf{u}}. \quad (52)$$

The alignment $\boldsymbol{\xi}^s \approx \mathbf{e}^s$ is fitted from the snapshots; it is not a kinematic identity. The iterative PROM-ANN therefore solves at fixed $\mathbf{e}$ for the state satisfying

$$\mathbf{r}(\boldsymbol{\xi}; \mathbf{e}) = \mathbf{B}_\xi^T \mathbf{R}(\tilde{\mathbf{u}}(\boldsymbol{\xi}); \mathbf{e}) = \mathbf{0}, \quad (53)$$

rather than imposing $\boldsymbol{\xi} = \mathbf{e}$. Because $\mathbf{B}_\xi$ depends on the state, the exact Newton derivative includes the curvature of the decoder manifold. Appendix D gives the coordinate construction, the reconstruction-error split, and this tangent explicitly.

For the comparison below, the affine and nonlinear models share the same retained span, $n_{\text{tot}} = n_{\text{tra}}$. The nonlinear model reduces the number of equilibrium unknowns from n_{tra} to $n = 3$; hyperreduction has not yet been introduced.

For compactness, the cubature and output formulas below retain $\mathbf{q}$ for the primary modal amplitudes; in the deployed nonlinear model, the decoder and adaptive-weight maps are evaluated through $\boldsymbol{\xi} = \mathbf{A}_m^{-1} \mathbf{q}$.

4.4 Hyperreduction of the projected equations

Although the PROMs reduce the number of equilibrium unknowns, their projected equations still require full-mesh element evaluations. Hyperreduction addresses this remaining cost by replacing each full sum with a sparse weighted sum. It is therefore a second approximation, introduced after the displacement representation has been chosen.

Two integration tasks must be distinguished: assembling the projected equilibrium equations and recovering the homogenized stress. Hernández et al. [39] formulate both tasks and discuss the use of either a common or a dedicated output rule. We use separate elementwise rules because their integrands and accuracy requirements differ. The construction is summarized here; Appendix E records the matrix conventions and fitting losses used in our implementation.

4.4.1 Fixed elementwise cubature

Let $\mathcal{E}$ be the microscopic element set and $\mathbf{L}_e$ extract element nodal values from $\mathbf{d}$. At fixed macroscopic input, define the reconstructed element state and its reduced tangent by

$$\mathbf{d}_e = \mathbf{L}_e(\mathbf{T}\tilde{\mathbf{u}} + \mathbf{g}_{\text{per}}), \quad \mathbf{B}_e = \mathbf{L}_e \mathbf{T} \mathbf{B}. \quad (54)$$

The full reduced residual is $\sum_{e \in \mathcal{E}} \mathbf{h}_e$, where $\mathbf{h}_e = \mathbf{B}_e^T \mathbf{f}_{\text{int},e}$. Fixed ECM approximates it as

$$\mathbf{r}_{\text{ECM}}(\mathbf{q}; \mathbf{e}) = \sum_{e \in \mathcal{Z}_r} w_e \mathbf{h}_e(\mathbf{q}; \mathbf{e}) = \mathbf{0}, \quad \mathcal{Z}_r \subset \mathcal{E}, \quad w_e \geq 0. \quad (55)$$

For the conventional HPROM, $\mathbf{B}$ is replaced by $\mathbf{V}_{\text{tra}}$ and the unknown is $\mathbf{q}_{\text{tra}}$. There is no ANN in this tier. A selected element retains its ordinary within-element quadrature. Consequently, selected elements, selected Gauss points, and independent coordinates are distinct counts.

Offline, the element contributions are collected over the training states and compressed by SVD. ECM then selects a sparse support and nonnegative multipliers that reproduce the retained integration modes [34]. If $\mathbf{G}_{\text{cub}}$ denotes those modes augmented by the constant-integration constraint and $\mathbf{b}_{\text{cub}} = \mathbf{G}_{\text{cub}} \mathbf{1}$, selection stops when

$$\frac{\|\mathbf{G}_{\text{cub}} \mathbf{w} - \mathbf{b}_{\text{cub}}\|_2}{\|\mathbf{b}_{\text{cub}}\|_2} \leq \varepsilon_{\text{ECM}}. \quad (56)$$

The tolerance ε_{SVD} controls the preceding integrand compression, whereas ε_{ECM} controls the sparse fit. For the residual rule, the element data are $\mathbf{h}_e$. For the output rule, they are the four components of the element contribution to the average first Piola stress; all four are retained because this tensor is not generally symmetric. Both rules impose $\sum_e w_e = N_{\text{el}}$, with physical quadrature measures already included in the element contributions. Their supports and weights remain fixed online.

4.4.2 Manifold-adaptive weights

A fixed rule must represent every sampled state with the same multipliers. MAW-ECM instead retains one common support but lets its weights vary over the reduced manifold [39]. Applied to the residual, this gives

$$\mathbf{r}_{\text{MAW}}(\mathbf{q}; \mathbf{e}) = \sum_{e \in \mathcal{Z}_r} w_e(\mathbf{q}) \mathbf{B}_e(\mathbf{q})^T \mathbf{f}_{\text{int},e}(\mathbf{d}_e) = \mathbf{0}. \quad (57)$$

Starting from a fixed ECM rule, the MAW-ECM pruning algorithm removes elements while redistributing nonnegative statewise weights; its local and graph-regularized stages are described in Ref. [39]. Pruning

identifies the support but not the continuous online map. We instantiate that final regression stage with neural fields. Because our equilibrium and stress rules are separate, we train two maps, $\mathcal{N}_r$ and $\mathcal{N}_s$, on independently pruned supports $\mathcal{Z}_r$ and $\mathcal{Z}_s$:

$$\begin{aligned} \mathbf{w}(\boldsymbol{\xi}) &= N_{\text{el}} \text{softmax}(\mathcal{N}_r(\boldsymbol{\xi}; \boldsymbol{\theta}_r)), \quad \sum_{e \in \mathcal{Z}_r} w_e = N_{\text{el}}, \\ \mathbf{v}(\boldsymbol{\xi}) &= N_{\text{el}} \text{softmax}(\mathcal{N}_s(\boldsymbol{\xi}; \boldsymbol{\theta}_s)), \quad \sum_{e \in \mathcal{Z}_s} v_e = N_{\text{el}}. \end{aligned} \quad (58)$$

Input standardization is absorbed into each network. The closure ANN $\mathcal{N}$ remains a different object: it predicts secondary displacement coordinates, not weights. The two weight fields are fitted against full-mesh element sums at each decoded training state. The scaled softmax guarantees positivity and normalization, but integration accuracy is learned. Thus the offline order is: train the closure, select and prune the two supports, and finally fit the two weight fields.

Online, the current reduced iterate is passed to both the closure and $\mathcal{N}_r$; $\mathcal{N}_s$ supplies the output weights. The supports and all network parameters remain fixed. Differentiating the adaptive residual produces three contributions: element stiffness, decoder curvature, and weight variation. We retain only the first in the microscopic modified-Newton iteration, while the consistent constitutive tangent differentiates the complete implemented residual and output, as described next. For the affine HPROM, the decoder curvature and weight-variation terms vanish, so its projected stiffness is the exact residual Jacobian.

4.5 Constitutive output

4.5.1 Equilibrium-based stress and consistent tangent

The reduced microscopic solver must return the same work-conjugate stress and tangent required by the FE² problem. The final stage of Figure 4 summarizes this common output path. Residual cubature alone is insufficient because equilibrium and effective stress involve different integrands. The separate output rule introduced above reconstructs the average first Piola stress as

$$\mathbf{P}_h(\mathbf{q}, \mathbf{e}) = \frac{1}{|Y|} \sum_{e \in \mathcal{Z}_s} v_e(\mathbf{q}) \sum_{g \in e} \omega_g \mathbf{P}_{\mu, g}(\mathbf{d}_e), \quad \mathbf{S}_h = \mathbf{F}^{-1} \mathbf{P}_h. \quad (59)$$

Here v_e is fixed for HPROM and state dependent for HPROM-ANN. Because the independently fitted output rule need not produce an exactly symmetric converted tensor, the implementation returns its symmetric work-conjugate components,

$$\mathbf{s}_h(\mathbf{q}, \mathbf{e}) = ((S_h)_{11}, (S_h)_{22}, \tfrac{1}{2}[(S_h)_{12} + (S_h)_{21}])^T, \quad (60)$$

and differentiates this actual output. The residual and stress supports may overlap; their union defines the element data required by the iterative constitutive evaluation, whereas only $\mathcal{Z}_r$ enters the equilibrium solve.

Reduced equilibrium defines an implicit state $\mathbf{q}^*(\mathbf{e})$ and hence the constitutive map $\mathbf{s}_h(\mathbf{q}^*(\mathbf{e}), \mathbf{e})$. If the residual Jacobian is nonsingular, differentiating equilibrium gives

$$\mathbf{D}_h = \mathbf{s}_{h, \mathbf{e}} - \mathbf{s}_{h, \mathbf{q}} \mathbf{r}_{h, \mathbf{q}}^{-1} \mathbf{r}_{h, \mathbf{e}}. \quad (61)$$

The partial derivatives hold the other argument fixed and are evaluated at the converged state. They include the periodic lifting, stress conversion, decoder curvature, and—for adaptive rules—weight variation. Equation (61), rather than the modified-Newton matrix used to reach microscopic equilibrium, supplies the tangent returned to the structural solver. Appendix E records how these derivatives are evaluated in the two implementations.

4.5.2 Direct decoder evaluation: D-HPROM-ANN

When input-informed coordinates closely approximate the loading parameters, the decoder can also be evaluated without solving reduced equilibrium, as discussed in the hyperelastic example of Ref. [39]. Here this defines a separate operating mode, D-HPROM-ANN, obtained by prescribing $\boldsymbol{\xi} = \mathbf{e}$ in Eq. (52):

$$\tilde{\mathbf{u}}_D(\mathbf{e}) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{A}_m \mathbf{e} + \overline{\mathbf{V}} \hat{\mathbf{N}}(\mathbf{e}). \quad (62)$$

The decoded state is inserted directly into the adaptive stress rule, using the same closure and stress-weight field as the iterative model. Only the stress support is needed; no residual cubature or reduced equilibrium

Jacobian is evaluated. The tangent is obtained by central differences of the complete direct stress map, updating the decoder, weights, lifting, and stress conversion at each perturbed input.

This route still evaluates the microscopic material law on selected elements; it is not a neural stress regressor or a learned macroscopic potential. The prefix D is our label for the direct operating mode, not a new projection method. Its accuracy depends on both the closure and the approximation $\xi \approx \mathbf{e}$. Conversely, solving projected equilibrium corrects the primary state only within the chosen manifold and cubature approximation, so it does not guarantee a smaller stress error than direct evaluation at every input.

4.6 Mechanical structure retained by the reduced models

Unlike a direct constitutive surrogate, every reduced model above continues to evaluate the original microscopic hyperelastic law; the iterative variants also enforce projected equilibrium. This preserves a direct micromechanical link, but it does not automatically make the deployed macroscopic stress the derivative of an effective potential.

For fixed weights and consistent kinematics, Galerkin internal forces can be derived from a weighted reduced energy, as in the structure-preservation results established for ECSW under their stated hypotheses [33]. Two features of the present deployment prevent us from transferring that conclusion without further assumptions. First, equilibrium and stress use independently fitted cubature rules and therefore need not be derivatives of the same energy. Second, when the weights depend on the reduced state, weighting energy derivatives is not equivalent to differentiating the weighted energy:

$$\nabla_{\mathbf{q}} \sum_e w_e(\mathbf{q}) \Pi_e(\mathbf{q}, \mathbf{e}) = \sum_e w_e(\mathbf{q}) \nabla_{\mathbf{q}} \Pi_e + \sum_e \Pi_e \nabla_{\mathbf{q}} w_e. \quad (63)$$

The final term is absent from Eq. (57), and nonnegative weights do not eliminate it. Consequently, we claim neither an exact effective potential nor a polyconvexity certificate for the deployed HPROM variants. Their consistent tangent is the derivative of the implemented stress map, not a certificate of conservativity or tangent symmetry.

4.7 Offline preparation and online evaluation

The accuracy–cost comparison in Section 6 concerns online constitutive evaluation. Before deployment, the offline stage collects HDM snapshots, constructs the reduced spaces and latent coordinates, fits the nonlinear closure, selects the cubature supports, trains adaptive weights where applicable, and builds the reduced meshes. These preparation costs are not included in the reported online timings.

To translate sparse cubature into wall-clock savings, each optimized reduced mesh contains only the selected elements and the nodes they require. The corresponding basis operators and periodic lifting are restricted to that support before online evaluation. Selected elements retain the original geometry, constituent law, and quadrature; their cubature weights approximate the omitted contributions. At each online query, the model reconstructs the required displacement entries, solves reduced equilibrium when its operating mode requires it, and returns effective stress and tangent.

Online cost therefore depends on the support size, the retained secondary dimension, the equilibrium iterations, and the derivative evaluation—not only on the number of primary coordinates. Table 2 closes the construction by comparing the three deployed evaluators with their HDM/FOM reference.

Table 2: Intrusive evaluators used in the structural accuracy–cost comparison and their HDM/FOM reference. All reduced models retain microscopic constitutive evaluations.

Model	State representation	Equilibrium unknowns	Integration	Learned component
HDM / FOM	$\mathbf{u} \in \mathbb{R}^N$	N	Full mesh	None
HPROM	$\tilde{\mathbf{u}} = \mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tra}} \mathbf{q}_{\text{tra}}$	n_{tra}	Fixed ECM	None
HPROM-ANN	$\tilde{\mathbf{u}}(\mathbf{q}) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \overline{\mathbf{V}} \mathcal{N}(\mathbf{q})$	n	Adaptive ECM	Secondary coordinates; residual and stress weights
D-HPROM-ANN	$\tilde{\mathbf{u}}_D(\mathbf{e}) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{A}_m \mathbf{e} + \overline{\mathbf{V}} \hat{\mathcal{N}}(\mathbf{e})$	None	Adaptive stress rule	Same closure and stress weights as HPROM-ANN

5 Constitutive assessment on periodic microstructures

In this section, we evaluate the proposed constitutive modeling approach on two periodic representative volume elements (RVEs) with different directional microstructures. Material A contains one elliptical cavity rotated relative to the cell axes; material B contains four elliptical cavities of different sizes and orientations. Both use the same isotropic matrix law and 20% porosity. Together they let us assess the learned energy in a cell with one prominent geometric direction and in a cell with several. We compare the surrogates with independent full-order RVE responses and, in material B, assess the effect of learning the paired-feature parameters against fixed-feature counterparts. The material-A models are subsequently deployed in the structural FE² calculation of Section 6. We distinguish the guarantees established by construction from the numerical accuracy and consistency checks reported here.

5.1 Microstructures and reference-model verification

5.1.1 Material A: single-cavity reference

The square periodic cell contains a centered elliptical void with area fraction 0.20, aspect ratio 2, and major-axis angle 30° relative to the cell axes; see Fig. 5. This orientation avoids the mirror-aligned configurations of the square cell. Centrosymmetry remains, but does not mean that the strain energy is an even function of strain.

The solid is compressible Neo-Hookean in plane strain. With $\mathbf{C}_\mu = \mathbf{F}_\mu^T \mathbf{F}_\mu$ and $J_\mu = \det \mathbf{F}_\mu > 0$, its local energy density is

$$\psi_\mu = \frac{\mu}{2}(\text{tr } \mathbf{C}_\mu - 2) - \mu \log J_\mu + \frac{\lambda}{2}(\log J_\mu)^2, \quad (64)$$

with Young’s modulus 1.628 GPa and Poisson ratio 0.4. Here μ and λ are the usual Lamé constants. This is the specific matrix law used by the reference solver, not a generic assertion about every Neo-Hookean variant. The production mesh contains 1546 quadratic triangles, 3240 nodes, and 4638 quadrature points. Periodicity and null-space constraints leave 6320 independent microscopic displacement coordinates. Cell averages use the complete reference cell measure.




Figure 5: Material A: cavity geometry (left) and working quadratic mesh (right). Dashed lines mark cavity major axes; signed angles are measured counterclockwise from X_1 . Only edges joining corner nodes are drawn, and ℓ denotes the cell side.

The production mesh was selected using a local stress-field criterion along a uniaxial-stress path, rather than only a homogenized average. On this path, E_{22} and $2E_{12}$ are solved so that $S_{22} = S_{12} = 0$. At $E_{11} = 0.20$, its maximum microscopic first Piola stress norm differs by approximately 0.486% from a 3095-element reference, whereas the average axial stress differs by approximately 0.001%. This is a mesh check at a specified state, not a uniform error bound over all multiaxial data.

5.1.2 Material B: multicavity reference

Material B retains the square periodic cell, the plane-strain Neo-Hookean matrix law and moduli, and the 20% total void fraction of material A. The void space is instead distributed among four ellipses with unequal sizes, aspect ratios, orientations, and off-center positions; see Fig. 6. Holding the matrix and porosity fixed isolates the change in cavity arrangement. Unlike the single rotated void in A, this cell presents several geometric directions within the same microstructure. We do not assign a material symmetry group to the resulting effective law.

The working mesh contains 4621 quadratic triangles and 9582 nodes. Periodic boundary conditions and full-cell averaging follow Section 2.




Figure 6: Material B: cavity geometry (left) and working quadratic mesh (right). Dashed lines mark cavity major axes; signed angles are measured counterclockwise from X_1 . Only edges joining corner nodes are drawn, and ℓ denotes the cell side.

At 64 prescribed multiaxial states, the working mesh was compared with an 8961-element mesh. Relative differences in homogenized energy, stress, and tangent remained below 0.1%; the sampled maximum microscopic first Piola stress norm differed by at most 4.06%. These finite comparisons support the working mesh, not a uniform convergence bound.

5.2 Data domains and learning protocol

Training domains and data roles. To define the training domain for material A, we first consider the problem it must solve. In Section 6, its learned law supplies the constitutive response throughout a coupon FE^2 calculation, where different macroscopic quadrature points visit different strains as loading proceeds. To locate those strains before training, we ran a displacement-controlled coupon prepass with a St Venant–Kirchhoff approximation based on the cell’s reference tangent. The prepass recorded the strain at every macroscopic quadrature point; it defined the sampling envelope, not the stresses used to train the models. Displacement control prevented errors in the approximate force response from determining that envelope. A second prepass with a softer response shifted the observed componentwise extrema by at most 5.1%; we expanded the spans by 40% to allow for such differences. Because the coupon is loaded in tension, we set the lower E_{11} bound to zero. In the engineering-strain coordinates $\mathbf{e} = (E_{11}, E_{22}, 2E_{12})^T$, the rounded bounds are

$$\mathbf{e}_{\min} = (0, -0.09865, -0.16807)^T, \quad \mathbf{e}_{\max} = (0.1955, 0.01635, 0.10793)^T. \quad (65)$$

We sample this box on a Cartesian grid, which also provides neighboring states for the reduced models of Section 4. Of its 4950 nonreference states, 4208 are used for fitting and 742 for validation (split seed 5). An independent set of 400 off-grid states tests interpolation.

Material B serves a different purpose: it tests the constitutive models on a cell with several geometric directions and permits a controlled comparison of fixed and learned features. It is not used in the coupon FE^2 calculation or a reduced model, so its sampling domain need not inherit A’s coupon-derived bounds or Cartesian grid. After full-order and mesh checks at multiaxial boundary states, we chose $E_{11}, E_{22} \in [-0.04, 0.20]$

and $2E_{12} \in [-0.08, 0.08]$. This box includes moderate normal compression as well as extension and shear within the checked range. Low-discrepancy samples in its interior and on its faces and corners supply 4200 fitting states; separate samplings supply 512 validation and 512 test states. Ten prescribed loading paths are reserved for evaluation. Figure 7 shows both sampling designs.




Figure 7: Strain-sampling domains for Material A (left) and Material B (right), in percent. Wireframes show the prescribed limits; colors identify fit, validation, and independent test. For legibility, each panel displays 240 fit, 110 validation, and 110 test coordinates selected deterministically; full counts appear in the text.

Optimization and model selection.

For material A, each network represents one scalar energy. Its first strain derivative supplies stress, and its second strain derivative at the undeformed state supplies the reference tangent (initial stiffness); these are not independent outputs. We fit all three quantities against RVE data. Stress has the largest weight because it is the response queried throughout the FE^2 calculation. The loss combines dimensionless relative mean-square errors for stress, energy, and reference tangent, weighted 1, 0.2, and 0.02, respectively. The tangent target comes from a separate coarse periodic RVE calculation.

Parameter fitting starts with 2600 full-batch Adam steps and ends with up to 400 L-BFGS iterations to refine the fit. Validation stress error determines which checkpoint is retained; the independent test and probe states are not consulted. The selected ICNN and ICKAN checkpoints both use random-initialization seed 16 and 32 paired features. Each has two hidden layers, with 24 nodes per layer for ICNN and 12 for ICKAN. Because the earlier baselines used different features, objectives, and training budgets, their comparison measures practical accuracy under those protocols, not compute-matched architectural superiority.

Before entering either neural core, each feature z_i is centered on its undeformed value z_i^0 and divided by a positive scale s_i . The center is updated as the learned feature powers change; the scale $s_i = \max(10^{-4}, \max_{\text{fit}} |z_i - z_i^0|)$ is set once from the initial spread over fitting states. This keeps input magnitudes comparable during optimization without changing the constitutive constraints. For the selected ICKAN, each connection learns coefficients for six integrated-hat basis functions; their knot centers are fixed from -1.2 to 1.2 at spacing $h = 0.48$.

Material B supplies the controlled feature comparison absent from A. We trained Free and the fixed- and learned-feature versions of ICNN and ICKAN at seeds 16, 29, and 47. Within each core, fixed and learned variants share the 32-feature initialization, data, and objective. Checkpoints were selected by validation stress error and locked before test and loading-path evaluation. All runs used validation-plateau learning-rate scheduling and stopping, then L-BFGS, with a 200,000-step Adam cap. Free reached this cap in every seed without a plateau; its Adam stage is therefore budget-limited.

5.3 Constitutive response and predictive accuracy

For material A, the selected ICNN and ICKAN models attain held-out aggregate stress errors of approximately 0.65%; Table 3 reports the independently audited selected checkpoints. Free has a corresponding test error

of 1.2643%. This comparison concerns the archived training protocols, not a matched-budget architecture study. Outside the training domain, Free has a lower aggregate probe error than either polyconvex model; Supplementary Section S1 preserves those results. Architectural admissibility does not guarantee extrapolation accuracy.

Table 3: Material A: independent-test constitutive errors of the selected polyconvex models, in percent.

Model	e_S^{test}	e_W^{test}	Worst test
ICNN	0.6491	0.2169	5.53
ICKAN	0.6551	0.2148	5.78

All values are percentages. Stress and energy errors are aggregate array norms; the last column is the maximum per-state stress error. Material A: 400 independent test states.

Draft task—evidence pending. Add FOM/model stress curves on predeclared axial, shear and combined loading paths, followed by the corresponding independent-test results for B. Inventory existing evaluations before requesting new FOM solves; do not connect unrelated scattered states as though they were a loading path. Report componentwise behavior and variability where needed to interpret aggregate errors.

5.4 Contribution of learned paired features

Draft task—evidence pending. Run a controlled comparison of fixed and learned admissible paired features, preferably for both A and B. Keep the core architecture, feature count, data, objective and optimization budget comparable; declare the fixed configuration and multiple initialization seeds before inspecting test errors. Retain model-specific reference normalization and growth terms. This comparison is not supplied by the archived baselines and has no results yet.

5.5 Mechanical consistency and scope

The following numerical checks concern the available material-A checkpoints. They complement, rather than replace, the construction guarantees in Section 3.

The archived independent audit records successful energy first- and second-derivative checks and saved-weight nonnegative-energy sufficient bounds for both selected models. The conservative finite-interval margins are 17.25 and 16.82 in normalized energy units. Random admissible-deformation tests support implementation consistency, but the polyconvexity argument is analytical. Neither aggregate error nor a successful bound on energy is a uniform bound on stress prediction error.

For the energy-based models, sampled rank-one curvatures provide an implementation check using the stress-dependent expression derived in Appendix A. The detailed sampling audit is retained in Supplementary Section S2. These finite samples do not establish polyconvexity; the ICNN and ICKAN guarantee follows from the analytical construction.

Closed-cycle work consistency follows from the scalar potential for Free, ICNN, and ICKAN. A numerical audit of the material-A constitutive outputs is reported in Supplementary Section S3.

Draft task—evidence pending. Repeat the reference-state, derivative, cycle and saved-parameter nonnegative-energy checks for the selected material-B models. Report bounds that fail or remain inconclusive. Existing A-checkpoint audits cannot certify B, and a second material trained separately does not demonstrate transfer between geometries.

6 Structural deployment for material A: accuracy and online cost

6.1 Purpose, boundary conditions, and coverage

We now use material A to compare learned constitutive laws and reduced microscopic solvers within the same nonlinear structural calculation. The coupon measures how their approximations affect structural accuracy and online cost. The dog-bone profile has overall length 165 mm, grip width 19 mm, narrow-section width 13 mm, narrow-section length 57 mm, and fillet radius 76 mm. It is a plane-strain computational benchmark using an ASTM-inspired in-plane profile, not a simulation certified to reproduce an ASTM tensile test.

The two end faces receive equal-and-opposite axial resultants of 100 kN per end, applied in 20 equal increments. Three point constraints remove rigid modes: the axial and transverse displacement at the left mid-edge node, and transverse displacement at the right mid-edge node. Entire end faces are not clamped. Consequently, a small skew of the deformed ends is not ruled out by the boundary conditions. The anisotropic response can couple axial extension to shear.

The macro mesh contains 690 quadratic triangles, 1501 nodes, and 2070 integration points. The reference extrusion thickness used to convert the two-dimensional stress integrals into the stated force resultants is 50 mm. This force convention is part of the plane-strain benchmark, not an ASTM specimen-thickness specification. The physical cell size is declared as 1.3 mm, corresponding to ten cells across the narrow width. This provides a scale-separation assumption, not a direct numerical validation of homogenization error at that ratio. All saved converged integration-point strains for every reported model remain inside the unchanged training box. This check does not automatically cover unsaved intermediate Newton iterates or certify another load case.




Figure 8: Nested FOM-FE² reference at 100 kN per end. Fields are displayed on the actual deformed mesh, without displacement amplification. Gauss-point strain and stress values are first averaged within each element and then patch-averaged to nodes for visualization only. These recovered contours are not the raw fields used in error norms and can smooth local extrema. Stresses are second Piola components. A three-dimensional von Mises stress is not reconstructed because the saved output does not include the necessary out-of-plane Cauchy stress.

6.2 Reduced-state representation for material A

The 4950 displacement snapshots provide the basis of the reduced microscopic models as well as the effective constitutive labels. Figure 9 separates POD compression from the strain-informed coordinate fit. The relative snapshot truncation norm for the retained 39-dimensional space is 9.94×10^{-7} . After its orthogonal rotation into primary and secondary spaces, $n = 3$ and $\bar{n} = 36$ in the notation of Section 4.

The fitted relation between the deployed primary coordinates and macroscopic strain has an aggregate relative array discrepancy of 3.60×10^{-5} on these snapshots; its maximum absolute component discrepancy is 2.55×10^{-5} . These quantities explain why macroscopic strain is a useful primary-coordinate predictor for this material. They do not imply exact equality at unsampled states, nor do they measure the closure error or the equilibrium residual. In particular, the direct substitution in D-HPROM-ANN is an additional approximation, even when the plot appears visually indistinguishable from the identity.

The secondary-coordinate closure is a 3–64–64–36 feed-forward network with two tanh hidden layers and a linear output. Its training script uses input standardization, the raw secondary-coordinate loss in Eq. (51), and a 90/10 internal fit/validation split with seed 20260903. Full-batch Adam has initial learning rate 2×10^{-3} , a validation-driven reduction schedule, and a maximum budget of 20000 epochs. The archived best relative validation mean-square error is 3.75×10^{-7} . This is an internal closure metric after a POD built from all 4950 snapshots, not an independent material-generalization estimate. Independent effective-stress and structural tests provide the latter assessment.

Each deployed adaptive cubature field has architecture 3–128–128–128–10, GELU hidden activations, and the scaled softmax output of Eq. (58). The residual and stress checkpoints use distinct ten-element supports. Their larger neural cores illustrate why the online cost cannot be inferred solely from the three equilibrium coordinates or from the element count. These reduction-training choices differ from those of the constitutive networks; the common material and data domain do not imply matched offline training budgets.




Figure 9: Diagnostics from the deployed 39-mode basis artifact. Left: relative POD snapshot truncation norm, before the primary/secondary rotation. Right: strain-informed primary coordinates versus the corresponding macroscopic strains; every tenth snapshot is plotted, whereas the reported discrepancies use all 4950 states. Neither diagnostic measures structural accuracy or proves validity outside the sampled domain.

6.3 Online models and implementation verification

The linear HPROM uses $n_{\text{tra}} = 39$ generalized coordinates, a residual rule supported on 135 elements, and a stress rule on 73 elements; their union contains 183 elements. The integrand-SVD tolerances are $\varepsilon_{\text{SVD}}^r = 10^{-3}$ and $\varepsilon_{\text{SVD}}^s = 10^{-5}$, retaining 134 and 72 integrand modes, respectively. The greedy integration tolerance in Eq. (56) is $\varepsilon_{\text{ECM}} = 10^{-6}$ for both rules, including the multiplier-sum constraint.

The iterative HPROM-ANN uses $n = 3$ primary and $\bar{n} = 36$ secondary coordinates, with ten residual and ten stress elements and a union of 19. Initial fixed-ECM candidates use $\varepsilon_{\text{SVD}} = 10^{-4}$ and $\varepsilon_{\text{ECM}} = 10^{-6}$; adaptive pruning targets ten elements per rule. Thus this final count is a prescribed pruning target assessed through integration and response errors, not a consequence of the greedy tolerance alone. The support-generation driver uses up to 500 sampled states and tests up to 20 removal candidates per elimination. Its positivity-enforcing phase is active from the first elimination; the available driver does not enable or supply the global graph coupling. The residual weight field is subsequently refitted on the same pruned support

with a longer training budget. Historical pruning diagnostics are not stored in the deployed checkpoints, so their filenames alone are not evidence of a graph-regularized run.

The direct model evaluates stress on ten elements and solves no microscopic equilibrium equations. All these counts refer to reduced computational meshes, not full-mesh evaluations followed by masking. Figure 10 shows their spatial distribution and the overlap of residual and output rules.




Figure 10: Element supports of the deployed hyperreduced microscopic models. The pale full-cell mesh provides location context only; discarded elements are not assembled online. Colors distinguish residual-only, stress-only, and shared elements. Supports are read from the same index arrays used to construct the reduced MDPA meshes.

For the optimized linear and direct implementations, the maximum normalized stress differences from the original implementations over the verification states are 4.71×10^{-14} and 4.20×10^{-13} . Corresponding tangent differences are 2.68×10^{-10} and 7.25×10^{-10} . Independent finite-difference checks give discrepancies of approximately 3.51×10^{-9} and 2.38×10^{-9} . Such comparisons check that optimization preserves the deployed model; they do not establish that either model exactly reproduces full microscopic equilibrium.

6.4 Error and timing definitions

For a saved nodal or integration-point array A , the reported relative error is

$$e_A = 100 \frac{\|\text{vec}(A - A_0)\|_2}{\|\text{vec}(A_0)\|_2}, \quad (66)$$

where A_0 is the nested reference on the same meshes. These are unweighted discrete array norms. The strain array stores engineering shear, whereas the stress array stores tensor shear. They are neither quadrature-weighted continuum L^2 norms nor invariant tensor Frobenius norms. The metric is kept identical across models; this does not make the three different field errors interchangeable.

Online wall time includes the complete macroscopic solve, microscopic or surrogate stress/tangent evaluations, communication, and lazy worker construction within the timed solve. Parent setup, offline data generation, training, reduced-mesh construction outside the timed region, and result export are excluded. The FOM uses 20 worker processes with one numerical-library thread per worker. The three optimized microscopic surrogates use the same process configuration and three sequential fresh complete runs. PANN evaluations are batched in one process and use the separately investigated thread settings shown in Table 4.

The macroscopic solve stops when the free-degree-of-freedom residual norm falls below 10^{-5} N or its ratio to the initial increment residual falls below 10^{-7} . The periodic FOM uses a relative microscopic displacement-increment tolerance of 10^{-10} , with continuation when needed. The archived FOM and selected PANN times each have one repetition in this comparison. Thus their ratios are observed online speedups under the stated implementations, not matched-thread efficiency estimates or repeated-run confidence intervals. Table 5 exposes the available reduced-model variability rather than selecting its fastest repetition. A publication-final timing study should repeat every tier under a common resource-budget protocol and retain all repetitions.

The affine HPROM uses its exact reduced residual Jacobian, whereas HPROM-ANN uses the projected-stiffness modified-Newton matrix described in Section 4.4.2; both use a microscopic increment tolerance

of 10^{-10} . Their effective constitutive tangents nevertheless account for the equilibrium sensitivity of their respective full residuals. The archived structural runs do not retain microscopic iteration totals. The reported speedup therefore compares these implemented workflows, including their different iteration matrices and derivative costs; it does not isolate a nonlinear-representation gain or quantify the iteration penalty of modified Newton.

Table 4: Complete coupon simulations: final-field errors against the nested FOM and archived online wall times.

Model	CPU	n_{run}	t [s]	t_0/t	e_u [%]	e_E [%]	e_S [%]
FOM	20p	1	4268.679	1.0	0.00000	0.00000	0.00000
HPROM	20p	3	67.348	63.4	0.00565	0.01653	0.00099
HPROM-ANN	20p	3	30.143	141.6	0.05200	0.14968	0.03727
D-HPROM-ANN	20p	3	3.788	1126.9	0.02324	0.04851	0.00967
ICNN	4t	1	1.734	2461.3	0.42468	0.77028	0.08880
ICKAN	8t	1	3.669	1163.5	0.38657	0.78932	0.07800
Free	12t	1	1.356	3146.9	0.92988	1.49464	0.13492
Regression	16t	1	0.946	4510.4	0.60620	0.81218	0.07378

p: worker processes, one numerical-library thread each; t: PyTorch threads. For $n_{\text{run}} = 3$, time is the median; for $n_{\text{run}} = 1$, it is the archived single run. These are online wall times, not offline-inclusive or matched-thread speedups.

Table 5: Sequential optimized-run repetitions used for median HPROM timings.

Model	t_1 [s]	t_2 [s]	t_3 [s]
HPROM	63.354	67.348	68.953
HPROM-ANN	27.947	34.628	30.143
D-HPROM-ANN	3.788	3.758	3.969

6.5 Accuracy–cost trade-off and Newton behavior

Figure 11 and Table 4 show distinct operating points rather than a single universally best model. The linear HPROM gives the smallest displacement error, approximately 0.00565%, at a median online time of 67.35 s. The iterative HPROM-ANN reduces this time to 30.14 s with a displacement error of approximately 0.0520%. The optimized iterative nonlinear model is therefore about 2.23 times faster than the *optimized* linear model; comparison against the old linear implementation would overstate this gain. The direct closure reduces the time to 3.79 s and has a displacement error of approximately 0.0232% on this path. Its lower error than the iterative nonlinear model is an observed outcome, not a consequence of a general error ordering.

The ICNN and ICKAN times are 1.73 s and 3.67 s, with displacement errors of approximately 0.425% and 0.387%. Their finite-strain structural accuracy is therefore consistent with a useful constitutive approximation, although it is poorer than the reduced microscopic models on this benchmark. Direct stress regression is fastest, and Free also converges. These baselines establish the accuracy and cost achievable on this loading path alongside the constrained models.

All models require four macroscopic iterations per increment in the archived runs. Consequently, this example provides no empirical evidence that polyconvexity reduced Newton iterations relative to Free or Regression. It establishes that the selected constrained models can be used successfully with the common structural solver, while maintaining their constitutive guarantees. A separate challenge problem would be needed to establish a solver-robustness advantage, and its reference solution would need to be validated independently.

7 Conclusions

Paired directional features with constrained trainable powers and determinant exponents provide a practical polyconvex constitutive construction for the anisotropic periodic material studied here. Their isotropic reference derivatives allow exact stress normalization through an affine determinant correction without tying the normalization coefficient to the volumetric barrier stiffness. The proof combines feature convexity in independent minors with core convexity and monotonicity, and includes every correction in the complete energy. Global nonnegative energy is supported separately by a sufficient bound evaluated for the selected parameters. Both a smooth ICNN and an integrated-hat spline-ICKAN achieve approximately 0.65% held-out aggregate stress error and below 0.43% displacement error in the nominal coupon simulation.




Figure 11: Online cost versus final displacement error on the common coupon problem. Horizontal segments show the observed range of three optimized HPROM runs; unbarred points are archived single PANN runs. Process and thread counts differ as detailed in Table 4. The plot does not include offline costs.

Optimized reduced microscopic models provide a complementary accuracy–cost reference: they are more accurate on this coupon, while direct learned constitutive evaluation is generally faster. All compared models converge with the same number of macroscopic iterations per increment. The structural results therefore demonstrate the practical use of the constrained energies alongside their analytical guarantees; they establish neither a Newton-convergence advantage nor a universal accuracy ranking.

A Rank-one paths and finite-strain energy curvature

This appendix supplies the algebra behind the two finite-strain curvature statements in criterion (C5). For a twice differentiable energy, nonnegative curvature along every rank-one perturbation is expressed by the Legendre–Hadamard condition

$$\partial_{\mathbf{F}\mathbf{F}}^2 W(\mathbf{F})[\mathbf{a} \otimes \mathbf{b}, \mathbf{a} \otimes \mathbf{b}] \geq 0 \quad \text{for all } \mathbf{a}, \mathbf{b} \in \mathbb{R}^2. \quad (67)$$

The argument below shows why the polyconvex representation in Eq. (17) implies this condition.

Determinant along a rank-one path. For arbitrary $\mathbf{F}, \mathbf{H} \in \mathbb{R}^{2 \times 2}$, expansion of the two products in the determinant gives

$$\begin{aligned} \det(\mathbf{F} + t\mathbf{H}) &= (F_{11} + tH_{11})(F_{22} + tH_{22}) - (F_{12} + tH_{12})(F_{21} + tH_{21}) \\ &= \det \mathbf{F} + t(F_{22}H_{11} + F_{11}H_{22} - F_{21}H_{12} - F_{12}H_{21}) + t^2 \det \mathbf{H}. \end{aligned} \quad (68)$$

The coefficient of t is $\text{cof } \mathbf{F} : \mathbf{H}$, since

$$\text{cof } \mathbf{F} = \begin{pmatrix} F_{22} & -F_{21} \\ -F_{12} & F_{11} \end{pmatrix}, \quad \mathbf{A} : \mathbf{B} = \sum_{i,j} A_{ij} B_{ij}. \quad (69)$$

For $\mathbf{H} = \mathbf{a} \otimes \mathbf{b}$, its entries are $H_{ij} = a_i b_j$, so

$$\det \mathbf{H} = a_1 b_1 a_2 b_2 - a_1 b_2 a_2 b_1 = a_1 a_2 (b_1 b_2 - b_2 b_1) = 0. \quad (70)$$

The determinant is therefore affine along a rank-one path:

$$\det(\mathbf{F} + t\mathbf{a} \otimes \mathbf{b}) = \det \mathbf{F} + t \text{cof } \mathbf{F} : (\mathbf{a} \otimes \mathbf{b}). \quad (71)$$

If $\mathbf{F}_B = \mathbf{F}_A + \mathbf{H}$ and $\mathbf{F}_s = (1-s)\mathbf{F}_A + s\mathbf{F}_B$, with $\mathbf{H} = \mathbf{a} \otimes \mathbf{b}$, then

$$(\mathbf{F}_s, \det \mathbf{F}_s) = (1-s)(\mathbf{F}_A, \det \mathbf{F}_A) + s(\mathbf{F}_B, \det \mathbf{F}_B). \quad (72)$$

For $0 \leq s \leq 1$ and positive endpoint determinants, the whole segment remains admissible. Convexity of $\mathcal{P}$ therefore yields

$$W(\mathbf{F}_s) \leq (1-s)W(\mathbf{F}_A) + sW(\mathbf{F}_B). \quad (73)$$

For a twice differentiable energy, the restriction to this line consequently has nonnegative second derivative, as required by Eq. (67).

Strain curvature and the chain rule. Let $\mathbf{H} \in \mathbb{R}^{2 \times 2}$ be an arbitrary perturbation and let t parametrize $\mathbf{F}(t) = \mathbf{F} + t\mathbf{H}$. Dots denote derivatives with respect to t , and $\text{sym}(\mathbf{A}) = (\mathbf{A} + \mathbf{A}^T)/2$ denotes the symmetric part of a matrix. Direct multiplication gives

$$\begin{aligned} \mathbf{E}(t) &= \frac{1}{2} [\mathbf{F}^T \mathbf{F} - \mathbf{I} + t(\mathbf{F}^T \mathbf{H} + \mathbf{H}^T \mathbf{F}) + t^2 \mathbf{H}^T \mathbf{H}] \\ &= \mathbf{E}(0) + t \text{sym}(\mathbf{F}^T \mathbf{H}) + \frac{1}{2} t^2 \mathbf{H}^T \mathbf{H}. \end{aligned} \quad (74)$$

Thus $\dot{\mathbf{E}}(0) = \text{sym}(\mathbf{F}^T \mathbf{H})$ and $\ddot{\mathbf{E}}(0) = \mathbf{H}^T \mathbf{H}$. Using the engineering-strain convention of Eq. (2), the chain and product rules give

$$\begin{aligned} \frac{dW}{dt} &= \mathbf{s} \cdot \dot{\mathbf{e}}, \\ \frac{d^2 W}{dt^2} &= \sum_{i,j} \frac{\partial s_i}{\partial e_j} \dot{e}_j \dot{e}_i + \sum_i s_i \ddot{e}_i = \dot{\mathbf{e}}^T \mathbf{D} \dot{\mathbf{e}} + \mathbf{S} : \ddot{\mathbf{E}}. \end{aligned} \quad (75)$$

Substituting $\ddot{\mathbf{E}}(0)$ gives

$$\left. \frac{d^2 W(\mathbf{F} + t\mathbf{H})}{dt^2} \right|_{t=0} = \dot{\mathbf{e}}(0)^T \mathbf{D} \dot{\mathbf{e}}(0) + \mathbf{S} : (\mathbf{H}^T \mathbf{H}). \quad (76)$$

Here $\mathbf{D} = \partial \mathbf{s} / \partial \mathbf{e}$ is the strain tangent and $\mathbf{S}$ is the second Piola–Kirchhoff stress, both evaluated at $t = 0$. The first term accounts for the change of stress with strain; the second accounts for the curvature of the strain path, weighted by the current stress. Even if $\mathbf{D}$ is positive semidefinite, the stress term can be negative in compression.

An analytical counterexample. Consider $W(\mathbf{E}) = \mu \mathbf{E} : \mathbf{E}$, where $\mu > 0$ is a constant with units of energy density. One has $\mathbf{S} = 2\mu \mathbf{E}$. Since

$$W(\mathbf{e}) = \mu(e_1^2 + e_2^2 + e_3^2/2), \quad \mathbf{D} = \text{diag}(2\mu, 2\mu, \mu), \quad (77)$$

the energy is strictly convex in engineering strain. At $\mathbf{F} = \lambda \mathbf{I}$, where $\lambda > 0$ is the common in-plane stretch, choose the rank-one perturbation $\mathbf{H} = \text{diag}(1, 0)$. Then

$$\dot{\mathbf{e}} = (\lambda, 0, 0)^T, \quad \mathbf{S} = \mu(\lambda^2 - 1)\mathbf{I}, \quad \mathbf{H}^T \mathbf{H} = \text{diag}(1, 0). \quad (78)$$

Hence the strain-tangent term is $2\mu\lambda^2$, and the stress term is $\mu(\lambda^2 - 1)$, giving

$$\left. \frac{d^2 W(\lambda \mathbf{I} + t\mathbf{H})}{dt^2} \right|_{t=0} = 2\mu\lambda^2 + \mu(\lambda^2 - 1) = \mu(3\lambda^2 - 1). \quad (79)$$

At $\lambda = 1/2$, the curvature is $-\mu/4$, violating Eq. (67) despite convexity in $\mathbf{E}$. This is an analytical counterexample, not a sampled observation about the RVE or a trained model. Compression alone does not force negative curvature: at $\lambda = 3/4$, the same expression is $11\mu/16 > 0$.

Equivalently, direct substitution gives

$$W(\lambda \mathbf{I} + t\mathbf{H}) = \frac{\mu}{4} [(\lambda + t)^2 - 1]^2 + \frac{\mu}{4} (\lambda^2 - 1)^2, \quad (80)$$

whose second derivative at zero is $\mu(3\lambda^2 - 1)$. The path is admissible near zero because $\det(\lambda \mathbf{I} + t\mathbf{H}) = \lambda(\lambda + t) > 0$ for $t > -\lambda$.

B Proofs for the paired features and complete energy

The proof follows the construction used in the main text. We first isolate one stretch-determinant contribution and derive the exponent bounds that make it convex. We then assemble two perpendicular contributions and compute their reference derivative, proving Proposition 1. Finally, we compose the resulting features with the monotone convex core and examine W_{learn} , W_{corr} , W_{vol} , and W_{dist} to establish Proposition 2.

B.1 Joint convexity of the directional building block

Every contribution in Eq. (24) has the form

$$f(\mathbf{x}, J) = \|\mathbf{x}\|^a J^{-b}, \quad \mathbf{x} = \mathbf{F}\mathbf{d}, \quad (81)$$

up to a positive constant, with $a = 2p$ or $2q$. It is therefore enough to determine when this building block is jointly convex in $(\mathbf{x}, J)$. Let $t = \|\mathbf{x}\| > 0$. Because f depends on $\mathbf{x}$ only through t , its Hessian separates into a block coupling the radial variable t to J and transverse directions orthogonal to $\mathbf{x}$. The radial-determinant block is

$$\begin{pmatrix} a(a-1)t^{a-2}J^{-b} & -abt^{a-1}J^{-b-1} \\ -abt^{a-1}J^{-b-1} & b(b+1)t^aJ^{-b-2} \end{pmatrix}. \quad (82)$$

For this symmetric 2×2 block to be positive semidefinite, its diagonal entries and determinant must be nonnegative. The diagonal entries are nonnegative for $a \geq 1$ and $b \geq 0$, while the determinant is

$$\begin{aligned} \det &= ab[(a-1)(b+1) - ab]t^{2a-2}J^{-2b-2} \\ &= ab(a-b-1)t^{2a-2}J^{-2b-2} \geq 0. \end{aligned} \quad (83)$$

The factor $a-b-1$ gives the upper bound $b \leq a-1$; if this bound is violated, the determinant is negative and the block has curvatures of opposite sign. The remaining, transverse Hessian eigenvalues are $at^{a-2}J^{-b} \geq 0$. Hence $a \geq 1$ and $0 \leq b \leq a-1$ make the complete Hessian positive semidefinite. Continuous extension gives convexity at $\mathbf{x} = 0$. For the constitutive evaluation on $\text{GL}^+(2)$, a nonzero reference direction cannot satisfy $\mathbf{F}\mathbf{d} = 0$, so this point does not introduce an interior singularity of the directional norm. Substituting $a = 2p$ gives $p \geq \frac{1}{2}$ and $0 \leq b \leq 2p-1$; the same argument with (q, c) establishes all bounds in Eq. (25).

B.2 Proof of Proposition 1

Proof. Convexity and objectivity. For any fixed learned parameter set, the directions $\mathbf{d}_i$ and $\mathbf{d}_i^\perp$ are independent of deformation. The maps $\mathbf{F} \mapsto \mathbf{F}\mathbf{d}_i$ and $\mathbf{F} \mapsto \mathbf{F}\mathbf{d}_i^\perp$ are therefore linear. Composing the convex building block above with these maps, multiplying by the positive factors $1/p_i$ and $1/q_i$, and adding the two contributions proves convexity of z_i in independent $(\mathbf{F}, J)$. A left proper rotation changes $\mathbf{F}\mathbf{d}$ to $\mathbf{Q}\mathbf{F}\mathbf{d}$ without changing its norm, so each contribution is objective.

Reference value and derivative. We now return to the physical relation $J = \det \mathbf{F}$, still holding every learned parameter fixed. At $\mathbf{E} = \mathbf{0}$, both directional invariants and J equal one, so the two contributions give $z_i^0 = 1/p_i + 1/q_i$. Moreover, $I_4(\mathbf{d}_i) = \mathbf{d}_i^T(\mathbf{I} + 2\mathbf{E})\mathbf{d}_i$ and $J^2 = (1 + 2E_{11})(1 + 2E_{22}) - 4E_{12}^2$ imply $\partial_{\mathbf{E}} I_4(\mathbf{d}_i) = 2\mathbf{d}_i \otimes \mathbf{d}_i$ and $\partial_{\mathbf{E}} J|_0 = \mathbf{I}$. Applying the product rule to the first contribution gives

$$\partial_{\mathbf{E}} \frac{I_4(\mathbf{d}_i)^{p_i}}{p_i J^{b_i}} \Big|_0 = \frac{1}{p_i} [2p_i \mathbf{d}_i \otimes \mathbf{d}_i - b_i \mathbf{I}] = 2\mathbf{d}_i \otimes \mathbf{d}_i - \frac{b_i}{p_i} \mathbf{I}. \quad (84)$$

The prefactor $1/p_i$ has cancelled the factor p_i produced by differentiating the stretch power. The perpendicular contribution has the same form with $(\mathbf{d}_i, p_i, b_i)$ replaced by $(\mathbf{d}_i^\perp, q_i, c_i)$. The two directional projectors satisfy $\mathbf{d}_i \otimes \mathbf{d}_i + \mathbf{d}_i^\perp \otimes \mathbf{d}_i^\perp = \mathbf{I}$; adding the contributions therefore removes the directional dependence at the reference and gives

$$\partial_{\mathbf{E}} z_i|_0 = 2(\mathbf{d}_i \otimes \mathbf{d}_i + \mathbf{d}_i^\perp \otimes \mathbf{d}_i^\perp) - \left(\frac{b_i}{p_i} + \frac{c_i}{q_i} \right) \mathbf{I} = \rho_i \mathbf{I}. \quad (85)$$

In the work-conjugate engineering-strain convention of Eq. (2), this tensor gradient corresponds to $\partial_{\mathbf{e}} z_i|_0 = \rho_i(1, 1, 0)^T$, proving Eq. (27). $\square$

Positive affine conditioning. The raw feature is converted to the core input by $y_i = (z_i - z_i^{\text{cen}})/s_i$, with $s_i > 0$ independent of deformation. Writing $\mathbf{X} = (\mathbf{F}, J)$,

$$\partial_{\mathbf{X}\mathbf{X}}^2 y_i = \frac{1}{s_i} \partial_{\mathbf{X}\mathbf{X}}^2 z_i, \quad y_i(\mathbf{X}_B) - y_i(\mathbf{X}_A) = \frac{z_i(\mathbf{X}_B) - z_i(\mathbf{X}_A)}{s_i}. \quad (86)$$

The first identity preserves positive-semidefinite curvature and the second preserves componentwise order. Thus the numerical conditioning applied before the monotone core changes neither property used in the composition proof below.

B.3 Proof of Proposition 2

Proof. Potential structure, objectivity, and stress symmetry (C1)–(C3). Every directional norm and J is unchanged by left proper rotations. The features can also be expressed through $\mathbf{d}_i^T \mathbf{C} \mathbf{d}_i$ and $\sqrt{\det \mathbf{C}}$. Hence the complete energy is objective, and differentiation with respect to symmetric strain gives a symmetric material stress. Differentiating this scalar energy supplies (C1) and the major symmetry of its strain tangent. Smoothness holds on $\text{GL}^+(2)$ because $J > 0$, nonzero material directions have nonzero deformed length, and both cores are C^2 . Because this argument uses left rotations only, it does not impose a material symmetry group.

Reference normalization (C4). Equation (37) shows that W_{learn} sets the reference energy, and Eq. (39) shows that W_{corr} cancels the learned reference stress. The calculations for W_{vol} and W_{dist} in Section 3.5 show that both growth terms have zero value and zero strain gradient at the reference. Thus the complete energy satisfies both reference conditions.

Polyconvexity (C5). Let $\mathbf{G} \in \mathbb{R}^{2 \times 2}$ and $j > 0$ be independent. Define $\mathcal{P}(\mathbf{G}, j)$ by the right-hand side of Eq. (36), replacing $\mathbf{z}$ by $\mathbf{z}(\mathbf{G}, j)$, J by j , and $\text{tr } \mathbf{C}$ by $\|\mathbf{G}\|_F^2$. For two such pairs $\mathbf{X}_1 = (\mathbf{G}_1, j_1)$, $\mathbf{X}_2 = (\mathbf{G}_2, j_2)$ and $0 \leq t \leq 1$,

$$\begin{aligned} H(\mathbf{z}(t\mathbf{X}_1 + (1-t)\mathbf{X}_2)) &\leq H(t\mathbf{z}(\mathbf{X}_1) + (1-t)\mathbf{z}(\mathbf{X}_2)) \\ &\leq tH(\mathbf{z}(\mathbf{X}_1)) + (1-t)H(\mathbf{z}(\mathbf{X}_2)). \end{aligned} \quad (87)$$

The first inequality uses feature convexity and core monotonicity; the second uses core convexity. Therefore W_{learn} is convex up to its constant reference subtraction. The correction $W_{\text{corr}} = -r(j-1)$ is affine for either sign of r , while W_{vol} and W_{dist} are convex in independent $(\mathbf{G}, j)$ because their nonaffine parts are positive multiples of $-\log j$, $(j-1)^2$, and $\|\mathbf{G}\|_F^2$. Thus $\mathcal{P}$ is convex and $W(\mathbf{F}) = \mathcal{P}(\mathbf{F}, \det \mathbf{F})$, establishing Eq. (17).

Growth (C6). Convexity of H gives its supporting-plane inequality at $\mathbf{z}^0$:

$$H(\mathbf{z}) \geq H(\mathbf{z}^0) + \sum_i h_i(z_i - z_i^0) \geq H(\mathbf{z}^0) - \sum_i h_i z_i^0. \quad (88)$$

The second inequality uses $h_i, z_i \geq 0$. Hence W_{learn} has a deformation-independent lower bound and cannot cancel the analytic growth terms. As $J \rightarrow 0^+$, the logarithmic part of $W_{\text{vol}} + W_{\text{dist}}$ diverges while W_{corr} remains bounded. As $J \rightarrow \infty$, the positive quadratic part of W_{vol} dominates every potentially negative linear or logarithmic term. Finally, if J remains in a compact subset of $(0, \infty)$ while $\|\mathbf{F}\|_F \rightarrow \infty$, the term $\varepsilon \|\mathbf{F}\|_F^2/2$ in W_{dist} diverges and all determinant-only terms remain bounded. These are the three stated growth limits. $\square$

C Sufficient lower bound for nonnegative energy

Define $k_i = 1/p_i + 1/q_i$ and $\eta_i = \rho_i/k_i$. Write $t = \|\mathbf{F} \mathbf{d}_i\|^2$, $u = \|\mathbf{F} \mathbf{d}_i^\perp\|^2$. Hadamard's determinant inequality gives $tu \geq J^2$. Since each feature is increasing in t, u , its minimum at fixed J occurs on $tu = J^2$. Minimizing

$$\frac{t^{p_i}}{p_i J^{b_i}} + \frac{J^{2q_i - c_i} t^{-q_i}}{q_i} \quad (89)$$

gives $t^{p_i+q_i} = J^{2q_i+b_i-c_i}$ and hence $z_i \geq k_i J^{\eta_i}$. This is a separate lower bound for each feature; simultaneous attainability for all directions is not assumed. It therefore remains valid after summing the supporting-plane inequalities.

Together with $\text{tr } \mathbf{C}/2 \geq J$ and the supporting plane of H , this feature bound gives

$$W \geq g(J) := \sum_i h_i k_i [J^{\eta_i} - 1 - \eta_i(J-1)] + (\alpha + \varepsilon)(J-1 - \log J) + \frac{\beta}{2}(J-1)^2. \quad (90)$$

Because $g(1) = g'(1) = 0$, it remains to establish $g''(J) \geq 0$ for every $J > 0$. Differentiation gives

$$J^2 g''(J) = \alpha + \varepsilon + \beta J^2 + \sum_i c_i^* J^{\eta_i}, \quad c_i^* = h_i k_i \eta_i (\eta_i - 1). \quad (91)$$

Only terms with $0 < \eta_i < 1$ can have negative coefficients. If $c_- = \sum_{c_i^* < 0} |c_i^*|$, the two inequalities

$$\alpha + \varepsilon + \sum_{\eta_i < 0} c_i^* - c_- \geq 0, \quad (92)$$

$$\beta + \sum_{\eta_i \geq 1} c_i^* - c_- \geq 0 \quad (93)$$

are sufficient on $0 < J \leq 1$ and $J \geq 1$, respectively. For the first region, $J^{\eta_i} \leq 1$ for negative-coefficient terms and $J^{\eta_i} \geq 1$ when $\eta_i < 0$. For the second region, divide $J^2 g''$ by J : all negative-coefficient powers then satisfy $J^{\eta_i-1} \leq 1$, while $J^{\eta_i-1} \geq 1$ for $\eta_i \geq 1$ and $\beta J \geq \beta$. Dropping the remaining nonnegative contributions gives the stated bounds. Their failure is inconclusive.

The selected parameters require a sharper check. The audit partitions $[10^{-4}, 10^4]$ into 4000 logarithmically spaced intervals. On each interval $[J_L, J_R]$, a positive power term is bounded below by the smaller endpoint value, whereas a negative-coefficient term is bounded below by $c_i^* J_R^{\eta_i}$, since its exponent lies in $(0, 1)$. The constant and quadratic terms are bounded below by $\alpha + \varepsilon$ and βJ_L^2 . Below the finite interval, the nonnegative negative-exponent terms increase and the magnitudes of the negative-coefficient terms decrease toward $J = 0$. Above it, factoring out J^{η_-} with $\eta_- = \max_{c_i^* < 0} \eta_i < 1$ yields a lower bound whose retained positive powers increase. These endpoint and tail inequalities cover every $J > 0$, not only the grid points. Positive margins establish the sufficient analytical argument for the saved parameters to the accuracy of its floating-point evaluation; random energy samples alone would not establish it.

D Latent-space closure details

This appendix supplies the algebra omitted from Section 4.3. Let $\Phi \in \mathbb{R}^{N \times n_{\text{tot}}}$ be the retained orthonormal POD basis, let $\mathbf{S}_{\text{snap}}$ collect the centered displacement snapshots, and collect the corresponding engineering-strain inputs in $\mathbf{M} = [\mathbf{e}^1, \dots, \mathbf{e}^{N_s}]$. Their POD coefficients are

$$\mathbf{Q}_{\text{POD}} = \Phi^T \mathbf{S}_{\text{snap}}. \quad (94)$$

The input-informed construction first identifies the least-squares map from these coefficients to strain,

$$\mathbf{T}_m = \mathbf{M} \mathbf{Q}_{\text{POD}}^\dagger \in \arg \min_H \|\mathbf{H} \mathbf{Q}_{\text{POD}} - \mathbf{M}\|_F^2, \quad (95)$$

where $\dagger$ denotes the Moore–Penrose pseudoinverse. Assuming rank three, the compact factorization

$$\Phi \mathbf{T}_m^T = \mathbf{V} \Sigma_m \mathbf{Z}_m^T, \quad \mathbf{A}_m = \Sigma_m^{-1} \mathbf{Z}_m^T \quad (96)$$

defines the three-column orthonormal primary block $\mathbf{V}$ and its coordinate scaling $\mathbf{A}_m$. The secondary block $\bar{\mathbf{V}}$ is an orthonormal complement within the retained POD span, ordered by snapshot energy. With $\mathbf{q} = \mathbf{A}_m \boldsymbol{\xi}$, the deployed decoder is

$$\tilde{\mathbf{u}}(\boldsymbol{\xi}) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{A}_m \boldsymbol{\xi} + \bar{\mathbf{V}} \hat{\mathcal{N}}(\boldsymbol{\xi}), \quad \boldsymbol{\xi}^s = \mathbf{T}_m \mathbf{Q}_{\text{POD}}^s. \quad (97)$$

Thus $\boldsymbol{\xi}^s \approx \mathbf{e}^s$ is a fitted snapshot relation, not a kinematic identity.

The displacement error separates the two approximation mechanisms. For $\mathbf{V}_{\text{tot}} = [\mathbf{V}, \bar{\mathbf{V}}]$ and projected coordinates $(\mathbf{q}^s, \bar{\mathbf{q}}^s)$,

$$\|\mathbf{u}^s - \tilde{\mathbf{u}}(\mathbf{q}^s)\|_2^2 = \|\bar{\mathbf{q}}^s - \mathcal{N}(\mathbf{q}^s)\|_2^2 + \|(I - \mathbf{V}_{\text{tot}} \mathbf{V}_{\text{tot}}^T)(\mathbf{u}^s - \mathbf{u}_{\text{ref}})\|_2^2. \quad (98)$$

The first term is the closure error within the retained span; the second is the POD truncation error. Increasing network capacity cannot recover a displacement component excluded from $\mathbf{V}_{\text{tot}}$.

Finally, let $\mathbf{B}_\xi = \partial_{\boldsymbol{\xi}} \tilde{\mathbf{u}}$, $\mathbf{R}$ be the constrained microscopic residual, and $\mathbf{K} = \partial_{\mathbf{u}} \mathbf{R}$. Differentiating the projected residual $\mathbf{r} = \mathbf{B}_\xi^T \mathbf{R}$ gives

$$\frac{\partial \mathbf{r}}{\partial \boldsymbol{\xi}} = \mathbf{B}_\xi^T \mathbf{K} \mathbf{B}_\xi + \sum_{j=1}^N R_j \frac{\partial^2 \tilde{u}_j}{\partial \boldsymbol{\xi} \partial \boldsymbol{\xi}}. \quad (99)$$

The second term is the curvature contribution of the decoder manifold. It need not vanish at reduced equilibrium because $\mathbf{B}_\xi^T \mathbf{R} = \mathbf{0}$ does not imply $\mathbf{R} = \mathbf{0}$.

E Elementwise cubature conventions

This appendix records the implementation conventions omitted from Section 4. The general ECM construction and MAW–ECM pruning algorithm follow Refs. [34, 39]; the expressions below specify how the deployed residual and stress rules are assembled.

Fixed rule. Let $\mathbf{c}_e^s \in \mathbb{R}^d$ contain the contribution of element e at training state s . Residual cubature uses $\mathbf{c}_e^s = \mathbf{h}_e^s$. Stress cubature uses the four components of

$$\mathbf{p}_e^s = \frac{1}{|Y|} \sum_{g \in e} \omega_g \mathbf{P}_{\mu,g}^s, \quad \mathbf{c}_e^s = \text{vec } \mathbf{p}_e^s = [p_{e,11}^s \quad p_{e,12}^s \quad p_{e,21}^s \quad p_{e,22}^s]^T. \quad (100)$$

One element occupies each row of the sampled integrand matrix,

$$\mathbf{M}_{\text{int}} = \begin{bmatrix} (\mathbf{c}_1^1)^T & \cdots & (\mathbf{c}_1^{N_c})^T \\ \vdots & & \vdots \\ (\mathbf{c}_{N_{\text{el}}}^1)^T & \cdots & (\mathbf{c}_{N_{\text{el}}}^{N_c})^T \end{bmatrix} \in \mathbb{R}^{N_{\text{el}} \times (dN_c)}. \quad (101)$$

A truncated SVD gives $\mathbf{M}_{\text{int}} \simeq \mathbf{U}_{\text{cub}} \mathbf{\Sigma}_{\text{cub}} \mathbf{V}_{\text{cub}}^T$. To enforce the multiplier sum without duplicating a constant component already represented by the retained integration modes, the fitting system uses

$$\mathbf{p}_{\perp} = (\mathbf{I} - \mathbf{U}_{\text{cub}} \mathbf{U}_{\text{cub}}^T) \mathbf{1}, \quad \mathbf{G}_{\text{cub}} = \begin{bmatrix} \mathbf{U}_{\text{cub}}^T \\ \mathbf{p}_{\perp}^T / \|\mathbf{p}_{\perp}\|_2 \end{bmatrix}, \quad \mathbf{b}_{\text{cub}} = \mathbf{G}_{\text{cub}} \mathbf{1}. \quad (102)$$

If the constant already lies in the retained span, the last row is redundant. Otherwise its exact reproduction gives $\sum_e w_e = N_{\text{el}}$. The reference quadrature measures are already contained in $\mathbf{c}_e^s$; the unknowns are element multipliers.

Adaptive neural fields. After MAW–ECM pruning selects a support $\mathcal{Z} = \{e_1, \dots, e_m\}$, neural fitting returns to the uncompressed contributions at each decoded state:

$$\mathbf{A}^s = \begin{bmatrix} \mathbf{c}_{e_1}^s & \cdots & \mathbf{c}_{e_m}^s \\ 1 & \cdots & 1 \end{bmatrix}, \quad \mathbf{b}^s = \begin{bmatrix} \sum_{e \in \mathcal{E}} \mathbf{c}_e^s \\ N_{\text{el}} \end{bmatrix}. \quad (103)$$

Thus the target uses the full-mesh sum, not the selected-column sum. Residual fitting uses $d = n$ and $\mathcal{Z}_r$; stress fitting uses $d = 4$ and $\mathcal{Z}_s$. With the scaled-softmax map of Eq. (58), the residual network minimizes

$$\mathcal{L}_{\text{cub}}(\boldsymbol{\theta}_r) = \frac{1}{|\mathcal{I}_{\text{fit}}|} \sum_{s \in \mathcal{I}_{\text{fit}}} \frac{\|\mathbf{A}^s \mathbf{w}(\boldsymbol{\xi}^s) - \mathbf{b}^s\|_2^2}{\max(\|\mathbf{b}^s\|_2^2, \delta_{\text{floor}})}, \quad (104)$$

and the stress network uses the analogous loss with $\mathbf{v}$. A preliminary fit to the sampled MAW–ECM weights initializes each network, but the final objective is integration accuracy; a reproducing weight vector is not unique.

Tangent evaluation. Both reduced implementations use Eq. (61), but evaluate its partial derivatives differently. For the affine HPROM, the assembled reduced stiffness and the analytic lifting derivative provide $d\mathbf{q}_{\text{tra}}^*/d\mathbf{e}$ through three linear right-hand sides; central differences are applied only to the stress along those linearized equilibrium directions. For HPROM–ANN, central differences at fixed $\mathbf{q}$ or fixed $\mathbf{e}$ evaluate all four partial-derivative blocks in Eq. (61), including decoder curvature and weight variation. The resulting small linear system does not require a new nonlinear microscopic solve for every perturbed strain. Numerical differentiation is an implementation choice; analytic or automatic differentiation would give the same formal tangent subject to their respective numerical tolerances.

References

- [1] C. Miehe, J. Schröder, J. Schotte, Computational homogenization analysis in finite plasticity simulation of texture development in polycrystalline materials, *Computer Methods in Applied Mechanics and Engineering* 171(3–4) (1999) 387–418. doi:10.1016/S0045-7825(98)00218-7.

- [2] V. G. Kouznetsova, W. A. M. Brekelmans, F. P. T. Baaijens, An approach to micro–macro modeling of heterogeneous materials, *Computational Mechanics* 27(1) (2001) 37–48. doi:[10.1007/s004660000212](https://doi.org/10.1007/s004660000212).
- [3] M. G. D. Geers, V. G. Kouznetsova, W. A. M. Brekelmans, Multi-scale computational homogenization: Trends and challenges, *Journal of Computational and Applied Mathematics* 234(7) (2010) 2175–2182. doi:[10.1016/j.cam.2009.08.077](https://doi.org/10.1016/j.cam.2009.08.077).
- [4] F. Feyel, A multilevel finite element method (FE²) to describe the response of highly non-linear structures using generalized continua, *Computer Methods in Applied Mechanics and Engineering* 192(28–30) (2003) 3233–3244. doi:[10.1016/S0045-7825\(03\)00348-7](https://doi.org/10.1016/S0045-7825(03)00348-7).
- [5] B. A. Le, J. Yvonnet, Q.-C. He, Computational homogenization of nonlinear elastic materials using neural networks, *International Journal for Numerical Methods in Engineering* 104(12) (2015) 1061–1084. doi:[10.1002/nme.4953](https://doi.org/10.1002/nme.4953).
- [6] P. Benner, S. Gugercin, K. E. Willcox, A survey of projection-based model reduction methods for parametric dynamical systems, *SIAM Review* 57(4) (2015) 483–531. doi:[10.1137/130932715](https://doi.org/10.1137/130932715).
- [7] F. Aldakheel, E. S. Elsayed, T. I. Zohdi, P. Wriggers, Efficient multiscale modeling of heterogeneous materials using deep neural networks, *Computational Mechanics* 72(1) (2023) 155–171.
- [8] K.A. Kalina, L. Linden, J. Brummund, M. Kästner, FEANN: an efficient data-driven multiscale approach based on physics-constrained neural networks and automated data mining, *Computational Mechanics* 71 (2023) 827–851. doi:[10.1007/s00466-022-02260-0](https://doi.org/10.1007/s00466-022-02260-0).
- [9] B. Liu, N. Kovachki, Z. Li, K. Azizzadenesheli, A. Anandkumar, A. M. Stuart, K. Bhattacharya, A learning-based multiscale method and its application to inelastic impact problems, *Journal of the Mechanics and Physics of Solids* 158 (2022) 104668.
- [10] A. Ghaderi, V. Morovati, R. Dargazany, A physics-informed assembly of feed-forward neural network engines to predict inelasticity in cross-linked polymers, *Polymers* 12(11) (2020) 2628. doi:[10.3390/polym12112628](https://doi.org/10.3390/polym12112628).
- [11] F. As’ad, C. Farhat, A staggered training framework for mechanics-informed neural networks in tractable multiscale homogenization with application to woven fabrics, *Computer Methods in Applied Mechanics and Engineering* 452 (2026) 118666.
- [12] F. As’ad, P. Avery, C. Farhat, A mechanics-informed artificial neural network approach in data-driven constitutive modeling, *International Journal for Numerical Methods in Engineering* 123 (2022) 2738–2759. doi:[10.1002/nme.6957](https://doi.org/10.1002/nme.6957).
- [13] D.K. Klein, M. Fernández, R.J. Martin, P. Neff, O. Weeger, Polyconvex anisotropic hyperelasticity with neural networks, *Journal of the Mechanics and Physics of Solids* 159 (2022) 104703. doi:[10.1016/j.jmps.2021.104703](https://doi.org/10.1016/j.jmps.2021.104703).
- [14] L. Linden, D.K. Klein, K.A. Kalina, J. Brummund, O. Weeger, M. Kästner, Neural networks meet hyperelasticity: A guide to enforcing physics, *Journal of the Mechanics and Physics of Solids* 179 (2023) 105363. doi:[10.1016/j.jmps.2023.105363](https://doi.org/10.1016/j.jmps.2023.105363). Consulted full text: arXiv:2302.02403v2.
- [15] D. Y. Gao, P. Neff, I. Roventa, C. Thiel, On the convexity of nonlinear elastic energies in the right Cauchy–Green tensor, *Journal of Elasticity* 127 (2017) 303–308. doi:[10.1007/s10659-016-9601-6](https://doi.org/10.1007/s10659-016-9601-6).
- [16] K. Xu, D. Z. Huang, E. Darve, Learning constitutive relations using symmetric positive definite neural networks, *Journal of Computational Physics* 428 (2021) 110072. doi:[10.1016/j.jcp.2020.110072](https://doi.org/10.1016/j.jcp.2020.110072).
- [17] J.M. Ball, Convexity conditions and existence theorems in nonlinear elasticity, *Archive for Rational Mechanics and Analysis* 63 (1977) 337–403.
- [18] P. G. Ciarlet, *Mathematical Elasticity, Volume I: Three-Dimensional Elasticity*, North-Holland, 1988.
- [19] J. Yvonnet, Q.-C. He, The reduced model multiscale method (R3M) for the non-linear homogenization of hyperelastic media at finite strains, *Journal of Computational Physics* 223 (2007) 341–368. doi:[10.1016/j.jcp.2006.09.019](https://doi.org/10.1016/j.jcp.2006.09.019).
- [20] J. A. Hernández, J. Oliver, A. E. Huespe, M. A. Caicedo, J. C. Cante, High-performance model reduction techniques in computational multiscale homogenization, *Computer Methods in Applied Mechanics and Engineering* 276 (2014) 149–189.

- [21] J. A. Hernández, A multiscale method for periodic structures using domain decomposition and ECM-hyperreduction, *Computer Methods in Applied Mechanics and Engineering* 368 (2020) 113192.
- [22] J. Barnett, C. Farhat, Y. Maday, Neural-network-augmented projection-based model order reduction for mitigating the Kolmogorov barrier to reducibility, *Journal of Computational Physics* 492 (2023) 112420. doi:10.1016/j.jcp.2023.112420.
- [23] N. Sibuet, S. Ares de Parga, J. R. Bravo, R. Rossi, Discrete physics-informed training for projection-based reduced-order models with neural networks, *Axioms* 14(5) (2025) 385. doi:10.3390/axioms14050385.
- [24] S. Ares De Parga, R. Tezaur, C.G. Hernández, C. Farhat, Nonlinear projection-based model order reduction with machine learning regression for closure error modeling in the latent space, *Computer Methods in Applied Mechanics and Engineering* 448 (2026) 118443. doi:10.1016/j.cma.2025.118443.
- [25] D. Amsallem, M. J. Zahar, C. Farhat, Nonlinear model order reduction based on local reduced-order bases, *International Journal for Numerical Methods in Engineering* 92(10) (2012) 891–916.
- [26] K. Lee, K. T. Carlberg, Model reduction of dynamical systems on nonlinear manifolds using deep convolutional autoencoders, *Journal of Computational Physics* 404 (2020) 108973.
- [27] J. Barnett, C. Farhat, Quadratic approximation manifold for mitigating the Kolmogorov barrier in nonlinear projection-based model order reduction, *Journal of Computational Physics* 464 (2022) 111348.
- [28] F. Fritzen, M. Hodapp, The finite element square reduced (FE^{2R}) method with GPU acceleration: towards three-dimensional two-scale simulations, *International Journal for Numerical Methods in Engineering* 107 (2016) 853–881. doi:10.1002/nme.5188.
- [29] S. Grimberg, C. Farhat, R. Tezaur, C. Bou-Mosleh, Mesh sampling and weighting for the hyperreduction of nonlinear Petrov–Galerkin reduced-order models with local reduced-order bases, *International Journal for Numerical Methods in Engineering* 122(7) (2021) 1846–1874.
- [30] S. Ares De Parga, J. R. Bravo, J. A. Hernández, R. Zorrilla, R. Rossi, Hyper-reduction for Petrov–Galerkin reduced order models, *Computer Methods in Applied Mechanics and Engineering* 416 (2023) 116298.
- [31] S. Chaturantabut and D. C. Sorensen. Nonlinear model reduction via discrete empirical interpolation. *SIAM Journal on Scientific Computing* 32(5) (2010) 2737–2764. doi:10.1137/090766498.
- [32] C. Farhat, P. Avery, T. Chapman, J. Cortial, Dimensional reduction of nonlinear finite element dynamic models with finite rotations and energy-based mesh sampling and weighting for computational efficiency, *International Journal for Numerical Methods in Engineering* 98(9) (2014) 625–662.
- [33] C. Farhat, T. Chapman, P. Avery, Structure-preserving, stability, and accuracy properties of the energy-conserving sampling and weighting method for the hyper reduction of nonlinear finite element dynamic models, *International Journal for Numerical Methods in Engineering* 102 (2015) 1077–1110. doi:10.1002/nme.4820.
- [34] J.A. Hernández, M.A. Caicedo, A. Ferrer, Dimensional hyper-reduction of nonlinear finite element models via empirical cubature, *Computer Methods in Applied Mechanics and Engineering* 313 (2017) 687–722. doi:10.1016/j.cma.2016.10.022.
- [35] A. T. Patera, M. Yano, An LP empirical quadrature procedure for parametrized functions, *C. R. Acad. Sci. Paris, Ser. I* 355 (2017) 1161–1167. doi:10.1016/j.crma.2017.10.020.
- [36] M. Yano, A. T. Patera, An LP empirical quadrature procedure for reduced basis treatment of parametrized nonlinear PDEs, *Computer Methods in Applied Mechanics and Engineering* 344 (2019) 1104–1123. doi:10.1016/j.cma.2018.02.028.
- [37] J. A. Hernández, J. R. Bravo, S. Ares de Parga, CECM: A continuous empirical cubature method with application to the dimensional hyperreduction of parameterized finite element models, *Computer Methods in Applied Mechanics and Engineering* 418 (2024) 116552.
- [38] J. R. Bravo, J. A. Hernández, S. Ares de Parga, R. Rossi, A subspace-adaptive weights cubature method with application to the local hyperreduction of parameterized finite element models, *International Journal for Numerical Methods in Engineering* 125(24) (2024) e7590.

- [39] J. A. Hernández, S. Ares de Parga, R. Rossi, Dimensional hyperreduction of nonlinear finite element models via empirical cubature with manifold-adaptive weights, arXiv:2609.03068v1 (2026). <https://arxiv.org/abs/2609.03068>
- [40] J. Schröder, P. Neff, Invariant formulation of hyperelastic transverse isotropy based on polyconvex free energy functions, *International Journal of Solids and Structures* 40(2) (2003) 401–445.
- [41] T. C. Gasser, R. W. Ogden, G. A. Holzapfel, Hyperelastic modelling of arterial layers with distributed collagen fibre orientations, *Journal of the Royal Society Interface* 3(6) (2006) 15–35. doi:10.1098/rsif.2005.0073.
- [42] B. Amos, L. Xu, J.Z. Kolter, Input convex neural networks, *Proceedings of Machine Learning Research* 70 (2017) 146–155. <https://proceedings.mlr.press/v70/amos17b.html>
- [43] K.A. Kalina, J. Brummund, W. Sun, M. Kästner, Neural networks meet anisotropic hyperelasticity: A framework based on generalized structure tensors and isotropic tensor functions, *Computer Methods in Applied Mechanics and Engineering* 437 (2025) 117725. doi:10.1016/j.cma.2024.117725. Consulted local version: arXiv:2410.03378v1 (2024).
- [44] D.K. Klein, K.A. Kalina, R. Ortigosa, J. Martínez-Frutos, M. Kästner, O. Weeger, Advances in polyconvex anisotropic hyperelasticity, arXiv:2605.27011v1 (2026). <https://arxiv.org/abs/2605.27011>
- [45] K. Linka, E. Kuhl, A new family of Constitutive Artificial Neural Networks towards automated model discovery, *Computer Methods in Applied Mechanics and Engineering* 403 (2023) 115731.
- [46] V. Tac, F. Sahli Costabal, A. Buganza Tepole, Data-driven tissue mechanics with polyconvex neural ordinary differential equations, *Computer Methods in Applied Mechanics and Engineering* 398 (2022) 115248. doi:10.1016/j.cma.2022.115248.
- [47] K.P. Abdolazizi, R.C. Aydin, C.J. Cyron, K. Linka, Constitutive Kolmogorov–Arnold networks (CKANs): Combining accuracy and interpretability in data-driven material modeling, *Journal of the Mechanics and Physics of Solids* 203 (2025) 106212. doi:10.1016/j.jmps.2025.106212.
- [48] P. Thakolkaran, Y. Guo, S. Saini, M. Peirlinck, B. Alheit, S. Kumar, Can KAN CANs? Input-convex Kolmogorov–Arnold Networks (KANs) as hyperelastic constitutive artificial neural networks (CANs), *Computer Methods in Applied Mechanics and Engineering* 443 (2025) 118089. doi:10.1016/j.cma.2025.118089.
- [49] V. Tac, K. Linka, F. Sahli-Costabal, E. Kuhl, A. Buganza Tepole, Benchmarking physics-informed frameworks for data-driven hyperelasticity, *Computational Mechanics* 73 (2024) 49–65. doi:10.1007/s00466-023-02355-2. Consulted local version: arXiv:2301.10714v1, with the title *Benchmarks for physics-informed data-driven hyperelasticity*.
- [50] M. R. Chmiel, J. L. Barnett, C. Farhat, Assessment of projection-based model order reduction for a benchmark hypersonic flow problem, *AIAA SciTech 2024 Forum* (2024), paper 2024-0250. doi:10.2514/6.2024-0250.